Correctness Conditioning as Selection in Chain-of-Thought Faithfulness Evaluation

Yueer Pan[1] **Yuexi Pan**[2]
[1]Zhejiang University
[2]University of Science and Technology of China
panjoel@zju.edu.cn panyuexi@mail.ustc.edu.cn

GitHub: CCS-CoTFE/CCS-CoTFE

Website: ccs-cotfe.github.io/CCS-CoTFE

ABSTRACT

Correctness conditioning changes the estimand in intervention-based chain-of-thought (CoT) faithfulness evaluation. We formalize this conditioning as a selection operator in a seven-variable structural causal model. The observed gap satisfies the exact identity $\Delta = \mathrm{Cov}(A, C)/P(C{=}1)$; under an explicit conditional-independence premise, the identity factors into difficulty and collider pathways. A structured audit of 18 published evaluations finds that 16 report a conditioning-dependent metric and that none analyses the effect as selection bias under the stated coding rule. On four sub-ceiling deployments, correctness conditioning raises measured faithfulness by $+7.4$ percentage points on the primary run (GLM-5.3, $n = 1000$, $p < 0.0001$); pooled Cochran–Mantel–Haenszel analysis gives OR $= 1.70$ ($p < 10^{-6}$). An assumption-free covariance decomposition assigns 42.8% of the GLM-5.3 gap to between-stratum variation in the interventionable-operation count, leaving the majority unattributed. When counterfactual compliance is observable for every instance, all-trace evaluation directly estimates population faithfulness without correctness filtering.

1 INTRODUCTION

A central question for trustworthy deployment of reasoning models is whether their chain-of-thought (CoT) traces faithfully reflect the computation producing the final answer [Lanham et al., 2023, Turpin et al., 2023]. A widely used approach is intervention-based evaluation: edit a reasoning step and measure whether the answer changes. Correctness conditioning changes the estimand of that evaluation; it is a selection step, not a neutral preprocessing choice. Our structured audit of 18 published evaluations (Section 4) finds 2 report filter-only results, 14 report both filtered and unfiltered results, and none analyses the selection effect under the coding rule we state.

We formalize this selection step with a seven-variable structural causal model (SCM, Figure 1) and derive the general bias formula $\Delta = \mathrm{Cov}(\eta, \pi)/P(C{=}1)$ under an explicit conditional-independence premise. The assumption-free identity $\Delta = \mathrm{Cov}(A, C)/P(C{=}1)$ is the estimand; the factorized form separates two candidate pathways without changing that estimand. We apply the framework at field level by auditing the literature and at deployment level by measuring the gap and evaluating adjustment by an observed difficulty proxy.

Two pathways are available in the SCM: a collider path opened by conditioning on correctness when latent shortcut competence exists ($R \to C \leftarrow S$), and difficulty confounding when task hardness jointly drives correctness and compliance ($D \to C$, $D \to A$). In our experiments only the second is active. The framework nonetheless derives the activation conditions for the first (an accuracy-axis hump at $P^* = \sqrt{p_s}$ and a ranking-reversal condition),

which we present as prospective thresholds: the reliance they assume is traced in one scoped setting (Appendix M), but the thresholds themselves are not.

Our contributions:

1. **A diagnostic framework for selection bias in CoT evaluation.** A seven-variable SCM formalizing the correctness filter as a selection operator, and a general bias formula $\Delta = \mathrm{Cov}(\eta, \pi)/P(C{=}1)$ that holds across pathways under a stated conditional-independence assumption (Section 3). The framework diagnoses correctness conditioning specifically as a sample-selection operator, complementing recent causal approaches to step-level faithfulness [Sun et al., 2026] and input-bias effects [Balasubramanian et al., 2025]. The formula identifies when filtering is harmless (uncorrelated propensities, no filtering, or constant compliance), turning "filter less" into a quantified condition.
2. **A field-level audit.** A structured audit of 18 published intervention-based CoT faithfulness evaluations finds that 16 report a conditioning-dependent metric and 2 report filter-only results with no unfiltered counterpart; none analyses the effect as selection bias under the stated coding rule (Section 4).
3. **Deployment-level measurement.** Applying the framework to four sub-ceiling deployments and two 100%-accuracy boundary controls measures a +7.4pp gap on GLM-5.3 ($p < 0.0001$; pooled OR 1.70). Operation-count strata account for 42.8% of that gap in an assumption-free covariance decomposition; the remaining component is unattributed (Section 6).

The framework also derives activation thresholds for collider bias via shortcut competence; neutral headline runs do not meet those thresholds, so the derivation is prospective. Exploratory activations under explicit-reliance framing use separate endpoints and are not pooled with the headline results (Appendix M). All-trace evaluation directly estimates population faithfulness whenever counterfactual compliance is observable on every instance (Section 3).

2 Related Work

Intervention-based CoT faithfulness. Chain-of-thought prompting [Wei et al., 2022] has become the standard reasoning interface for large language models. Lanham et al. [2023] introduce behavioral interventions (truncation, paraphrase, mistake insertion) and partially filter by correctness. Turpin et al. [2023] show CoT can be biased by injected cues while preserving plausible reasoning. Bentham et al. [2024] is the closest qualitative antecedent to our concern: they show that "unfaithfulness" measures can be confounded by answer accuracy, but they frame this as a metric-degeneracy problem rather than a formal selection analysis, and restrict the diagnostic to a single setting. More recent work frames faithfulness through information flow [Jia et al., 2026], counterfactual simulation [Hase and Potts, 2026], step-level importance [Bogdan et al., 2025], unlearning [Tutek et al., 2025], and benchmarking [Shen et al., 2026].

Two 2026 works diagnose a faithfulness-measurement failure but not the one we study. Li et al. [2026] measure step-level alignment between latent answer commitment and the explicit trace (61.9% average alignment); their object is the trace's internal timing, ours the estimate computed over a filtered subset of traces, and the failure modes are independent; a trace can be step-level faithful yet yield a biased aggregate. Young [2026] show classifier choice alone reverses rankings (regex 74.4% vs. LLM judge 69.7% on identical data); their distortion

enters after the sample is fixed, ours before any instrument runs, through which instances survive the correctness filter. Adopting one classifier does not remove our bias and reporting all traces does not remove theirs. Our contribution is orthogonal: correctness conditioning introduces systematic measurement bias, and we derive a closed-form characterization of it.

Recent work applies causal reasoning to CoT evaluation from complementary angles. Sun et al. [2026] use controlled perturbations as instrumental signals to assess step-level faithfulness within a single chain (intra-chain dependence), while Fu et al. [2025] model CoT steps as causal variables to ensure both correctness and interpretability. Balasubramanian et al. [2025] study how input-level biases (text and image) affect CoT faithfulness articulation in vision-language models. Pathak et al. [2026] propose confidence-based trace filtering as an alternative evaluation strategy. Our framework addresses a distinct and complementary question: how the common practice of conditioning the evaluation *sample* on answer correctness introduces selection bias into aggregate faithfulness measures, a sample-level effect that operates regardless of step-level evaluation quality.

Selection bias in evaluation. Evaluation biases in LLM systems have been studied for position effects [Shi et al., 2025] and metric disagreement. McMilin [2022] identifies collider bias in LLM training data curation; we address a distinct manifestation, evaluation-time correctness filtering. Zaman and Srivastava [2025] analyze faithfulness metrics through a causal lens, but their focus is metric validity whereas we analyze the sample-selection operator. Gur-Arieh et al. [2026] meta-evaluate metrics against ground truth; they evaluate metric *accuracy*, we evaluate metric *bias* from sample selection. Our formula is an instance of Heckman-type sample-selection [Heckman, 1979] and of selection bias from conditioning on a collider [Elwert and Winship, 2014]; the conditioning-versus-truncation distinction maps onto the missing-data mechanisms of Rubin [1976], since correctness filtering discards instances missing not at random [Rubin, 1976, Mohan and Pearl, 2021]. Adjacent work studies selection under other guises: selective classification with abstention [El-Yaniv and Wiener, 2010, Geifman and El-Yaniv, 2017], the selective-labels problem [Lakkaraju et al., 2017], and graphical recoverability criteria [Bareinboim and Pearl, 2012, Bareinboim et al., 2014]. Those works treat selection as the modelling target; here it is an unrecognized by-product of a standard protocol, which is why we frame the result as a diagnostic rather than a new estimator. Parcalabescu and Frank [2024] separate faithfulness from self-consistency and Atanasova et al. [2023] test explanation faithfulness; both target the metric's semantics, whereas we target the sample it is computed on.

Causal inference and shortcut learning. Our SCM follows Pearl's structural approach [Pearl, 2009], and the intervention analysis builds on causal-mediation methods for neural models [Vig et al., 2020]; the collider $R \rightarrow C \leftarrow S$ is an instance of Berkson's paradox [Berkson, 1946]. Conditioning on post-treatment variables is a well-known bias source in observational studies [Rosenbaum, 2002], and our variable S maps to the shortcut taxonomy of Geirhos et al. [2020]. Recent evidence that LLMs use hidden cues [Arcuschin et al., 2025, Chen et al., 2025] and that monitoring promotes obfuscation [Baker et al., 2025] makes our bias diagnosis practically relevant.

3 A DIAGNOSTIC FRAMEWORK FOR SELECTION BIAS

We build a formal diagnostic framework for the selection step that correctness filtering performs. We define a structural causal model (SCM) for intervention-based evaluation, derive the covariant bias form that holds whichever named pathway carries the dependence, and apply it to the two pathways the SCM admits: difficulty confounding, and collider bias via shortcut competence, whose activation conditions we derive.

3.1 Structural Causal Model

Definition 1 (Faithfulness Evaluation SCM)**.** *Let* $(\mathbf{U}, \mathbf{V}, \mathcal{F})$ *be a stochastic SCM with endogenous variables* $\mathbf{V} = \{D, S, R, T, C, I, A\}$ *and mutually independent exogenous noise* $\mathbf{U} = \{U_D, U_S, U_R, U_T, U_C, U_I, U_A\}$*. The Bernoulli kernels below are shorthand for structural functions of the corresponding independent noise terms:*

$$D = f_D(U_D), \tag{1}$$
$$S \sim \text{Bernoulli}(p_S(D)), \tag{2}$$
$$R \sim \text{Bernoulli}(p_R(D)), \tag{3}$$
$$T \sim \text{Bernoulli}(p_T(R, D)), \tag{4}$$
$$C \sim \text{Bernoulli}(\pi(R, S, T, D)), \tag{5}$$
$$I \sim \text{Bernoulli}(\rho), \tag{6}$$
$$A \sim \text{Bernoulli}(\eta(R, S, T, D, I)), \tag{7}$$

where D *is task difficulty,* S *is shortcut competence (whether the model has a non-chain-of-thought route to the answer),* R *is chain-of-thought (CoT) reliance (whether the model genuinely uses the visible reasoning trace),* T *is trace correctness,* C *is answer correctness (the collider),* I *is randomized intervention assignment, and* A *is post-intervention answer change (the compliance outcome).*

The Bernoulli parameterization is the natural exponential-family form for binary outcomes. In a typical pipeline, $C = 1$ corresponds to filtering for traces the model originally answered correctly, and A indicates whether the answer changes appropriately after an intervention on its reasoning chain. Reporting faithfulness only over correct traces is therefore a selection operator: it conditions on $C = 1$. Two structures in the SCM let that conditioning bias the estimate:

$$\underbrace{R \to C \leftarrow S}_{\text{collider (needs } S>0)} \qquad \text{and} \qquad \underbrace{D \to C,\ D \to A}_{\text{difficulty confounding}}. \tag{8}$$

Section 3.3 treats the confounding structure, which our covariate explains; Section 3.4 treats the collider, whose activation conditions we derive but which our models do not satisfy (note $T \perp S \mid R, D$).

3.2 The General Bias Formula

Within a fixed difficulty stratum $D = d$, define the T-marginalized cell probabilities:

$$\pi_{rs} \coloneqq P(C = 1 \mid R = r, S = s, D = d), \tag{9}$$
$$\eta_{rs} \coloneqq P(A = 1 \mid R = r, S = s, I = 1, D = d). \tag{10}$$

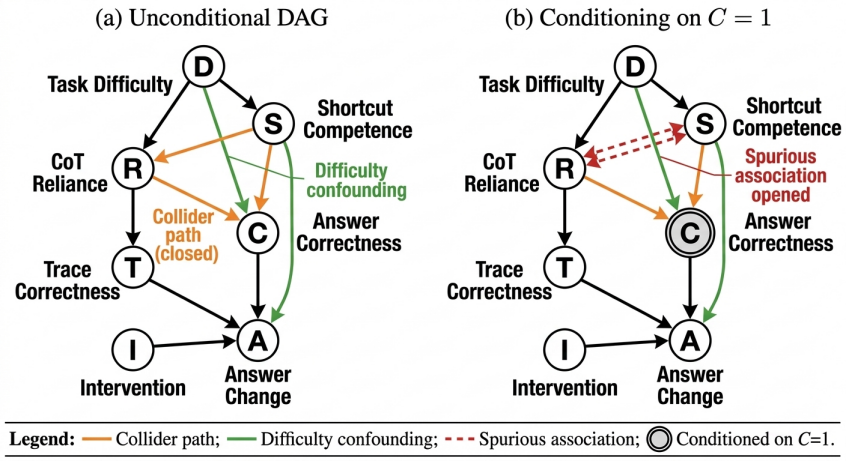


Figure 1: Structural causal model for faithfulness evaluation. (a) The unconditional DAG shows the collider structure $R \to C \leftarrow S$ (orange) and difficulty confounding pathways $D \to C$, $D \to A$ (green). (b) Conditioning on $C = 1$ (double border) opens a spurious association between CoT reliance R and shortcut competence S (dashed red), biasing the faithfulness estimate.

Assume $R \perp S \mid D = d$ with $P(R = 1 \mid D = d) = p_r$ and $P(S = 1 \mid D = d) = p_s$. In our SCM, $R = f_R(D, U_R)$ and $S = f_S(D, U_S)$ where U_R and U_S are independent exogenous variables. By the Markov property, $R \perp S \mid D$ follows from the SCM structure; reasoning ability and shortcut competence share only the common cause of task difficulty.

Definition 2 (Population and Selected-Trace Faithfulness)**.** *For causal interventions ($I = 1$), within a fixed difficulty stratum $D = d$:*

$$F_{\text{pop}}^{(d)} := P(A = 1 \mid I = 1, D = d), \tag{11}$$

$$F_{\text{sel}}^{(d)} := P(A = 1 \mid I = 1, C = 1, D = d), \tag{12}$$

$$\Delta^{(d)} := F_{\text{sel}}^{(d)} - F_{\text{pop}}^{(d)}. \tag{13}$$

The marginal analogues reported in Table 1 drop the stratification, with F_{pop}, F_{sel}, and $\Delta := F_{\text{sel}} - F_{\text{pop}}$ defined by the same expressions without conditioning on D.

Proposition 1 (General Bias Formula)**.** *Assume $A \perp C \mid \{R, S, I, D\}$ (compliance is conditionally independent of correctness given the strategy type and intervention). Then the selection bias equals the covariance of intervention-response propensity and correctness propensity, normalized by the selection probability:*

$$\Delta = \frac{\text{Cov}(\eta_{RS}, \pi_{RS} \mid D = d)}{P(C = 1 \mid D = d)}, \tag{14}$$

where η_{RS} and π_{RS} are random cell-level functions of (R, S) under the population distribution.

Proof. Expanding F_{sel} using Bayes' rule:

$$F_{\text{sel}} = \sum_{r,s} \eta_{rs} \cdot \frac{\pi_{rs}\, P(R{=}r)\, P(S{=}s)}{P(C{=}1)} = \frac{E[\eta_{RS}\, \pi_{RS}]}{P(C{=}1)}. \tag{15}$$

Since $F_{\text{pop}} = E[\eta_{RS}]$ and $P(C{=}1) = E[\pi_{RS}]$:

$$\Delta = \frac{E[\eta_{RS}\,\pi_{RS}]}{E[\pi_{RS}]} - E[\eta_{RS}] = \frac{E[\eta_{RS}\,\pi_{RS}] - E[\eta_{RS}]\,E[\pi_{RS}]}{E[\pi_{RS}]} = \frac{\text{Cov}(\eta_{RS}, \pi_{RS})}{P(C{=}1)}. \quad (16)$$

□

What is assumed and what is not. The estimand is assumption-free, and the two-line derivation shows why. With $F_{\text{sel}} = E[AC]/E[C]$ and $F_{\text{pop}} = E[A]$,

$$\Delta = \frac{E[AC]}{E[C]} - E[A] = \frac{E[AC] - E[A]\,E[C]}{E[C]} = \frac{\text{Cov}(A, C)}{P(C{=}1)}, \quad (17)$$

using $P(C{=}1) = E[C]$. Equation 17 is an algebraic identity on the joint distribution of the observed pair (A, C): no latent variable, no independence premise. It is the object we estimate throughout. The assumption enters only in *factoring* it: in equation 5–equation 7, trace correctness T is a direct argument of both C and A, so conditioning on $\{R, S, I, D\}$ does *not* block $C \leftarrow T \rightarrow A$ and the premise $A \perp C \mid \{R, S, I, D\}$ is imposed rather than entailed. It is what licenses reading equation 14 as attributing the bias to (η_{RS}, π_{RS}); the two covariances are equal only when it holds, and $\text{Cov}(\eta_{RS}, \pi_{RS})$ is not the same object as $\text{Cov}(A, C)$.

The same factoring can be done assumption-free by stratifying on an *observed* variable. With D the measured difficulty proxy, the standard covariance decomposition gives

$$\text{Cov}(A, C) = \underbrace{\text{Cov}(E[A \mid D],\ E[C \mid D])}_{\text{between-stratum}} + \underbrace{E[\text{Cov}(A, C \mid D)]}_{\text{within-stratum}}, \quad (18)$$

which holds for any joint distribution of (A, C, D). We use equation 18 in Section 6 to measure the difficulty contribution directly; its within-stratum term is *unattributed*, not evidence of a collider.

equation 14 is the framework's core result: *whichever* named pathway carries the dependence, the bias takes the covariance form. Correctness filtering is unbiased only when $\text{Cov}(A, C) = 0$, when $P(C{=}1) = 1$, or, under the premise, when η_{RS} is constant across cells. Read as an attribution, equation 14 requires the premise above; read as a description of Δ itself, equation 17 does not. The theory identifies the stratum-specific $\Delta^{(d)}$; the marginal Δ of Table 1 is a weighted mixture of those and is treated as a distinct estimand (Appendix D).

3.3 Difficulty Confounding

When task difficulty D jointly influences correctness ($D \rightarrow C$) and compliance ($D \rightarrow A$), conditioning on $C = 1$ selects easier tasks where compliance is higher. This is the pathway our covariate measures (Section 6). We report an exact between-stratum share for one observed proxy (operation count), not a causal contribution of latent difficulty: the share is 21.9–42.8% and the remainder is within-stratum and unattributed. Identifying a pathway would require intervening on the confounder itself, randomizing difficulty while holding the other causes of correctness fixed; stratification does not reach that bar.

Proposition 2 (Difficulty-Confounding Bias)**.** *Let* $\pi(d) = P(C = 1 \mid D = d)$ *and* $\eta(d) = P(A = 1 \mid I = 1, D = d)$. *When both* π *and* η *decrease with difficulty (harder*

tasks have lower accuracy and lower compliance), conditioning on $C = 1$ selects easier instances and inflates the faithfulness estimate:

$$\Delta = \frac{\mathrm{Cov}(\eta(D), \pi(D))}{P(C = 1)}, \tag{19}$$

the same covariance form as Proposition 1 with (R, S) replaced by the observed variable D. Whenever harder tasks produce both lower accuracy and lower compliance ($\mathrm{Cov}(\eta(D), \pi(D)) > 0$), the bias is strictly positive.

Proposition 2 holds under the additional assumption $A \perp C \mid \{D, I\}$, satisfied when compliance depends only on task difficulty and intervention type. As above this premise is *imposed, not derived*: D is a direct argument of both C and A in equation 5–equation 7, so it is residual dependence beyond D that the assumption excludes. Its practical value over the collider form is that D is *measured* per instance, so the covariance can be estimated and adjusted for directly, the basis of the correction in Section 6. Because both pathways are instances of equation 14, all-trace evaluation removes the covariance term regardless of which is operating.

3.4 Prospective Analysis: Shortcut Competence and Collider Bias

A second pathway is available in the SCM: when latent shortcut competence $S > 0$ exists, conditioning on correctness opens a collider ($R \to C \leftarrow S$), amplifying the bias beyond what difficulty confounding alone produces. Under a strict-reliance special case, we derive a closed-form gap (Proposition 3), show it is monotonically decreasing in shortcut prevalence (Corollary 2), hump-shaped in CoT reliance with peak at $p_r^* = \sqrt{p_s}/(1 + \sqrt{p_s})$ (Corollary 3), and identify conditions under which two models' faithfulness rankings reverse (Proposition 4). These results require an active shortcut branch, which under neutral framing none of our deployments exhibits (§6: 0% measured shortcut use), though multiple deployments under an explicit-reliance framing do (Appendix M), unpooled with the headline runs. The full derivations, proofs, and activation thresholds are in Appendix F; we present them as prospective.

3.5 Implications for Evaluation Practice

equation 14 has three consequences: the bias vanishes only under restrictive conditions (uncorrelated compliance and correctness, perfect accuracy, or constant compliance); it peaks at intermediate CoT reliance (Corollary 3), where evaluation is most valuable; and evaluating on all traces computes F_{pop} directly, eliminating the covariance term. The last point carries a scope condition: all-trace evaluation requires a compliance signal computable on every instance, including the incorrect ones, so where no counterfactual oracle exists—most natural-language settings—the practical fallback is to report both F_{sel} and F_{pop}.

4 Literature Audit

The framework in Section 3 identifies a bias source; whether it matters depends on how widely the triggering practice is used. We audited intervention-based CoT faithfulness evaluations published between 2023 and 2026, recording for each whether the reported metric conditions on answer correctness and whether the paper analyses the resulting selection

effect. Table 4 (Appendix H) gives the full record. The working coding ledger is retained by the authors and is not part of the public release.

Retrieval protocol. Candidates were gathered from arXiv (cs.CL and cs.AI listings for 2023–2026 plus keyword sweeps for *chain-of-thought faithfulness*, *unfaithful reasoning*, and *intervention*), the ACL Anthology, and the reference lists and citation graph of the papers already included; NeurIPS/ICML/ICLR proceedings were checked for conference versions of the same work. The window is publication dates 2023-01 to 2026-09. We screened titles and abstracts first, then full texts; the seven candidates excluded on the inclusion rule were retained in the authors' audit notes. This is a structured search, not a systematic review with a registered protocol, so we report it as a described procedure rather than claiming exhaustive coverage.

Inclusion and classification. A paper enters if it intervenes on a reasoning trace and reports a faithfulness-style metric; survey, purely mechanistic, and accuracy-meta-analysis papers are excluded, yielding 18 papers. We record what the *reported metric* conditions on, not what the authors intend: **Explicit** (n=2) reports *only* correctness-filtered results; **Partial** (n=14) filters some analyses but also reports the unfiltered number, so at least one headline figure is conditioning-dependent but a reader can recover from it; **None** (n=2) avoids filtering. No audited paper fell in the implicit category.

Two counts that should not be conflated. 2 of 18 papers report filter-only results: their published number carries the bias with no unfiltered counterpart a reader could recover from, the strictest reading. Counting also the 14 that report both gives **16 of 18** whose reported metrics are conditioning-dependent, leaving 2 that avoid filtering by design. **Conditional on the coding rule below, 0 of the 18 analyse the effect as a selection bias.** We mark that zero as conditional deliberately: it is the rule's verdict, and a looser rule, one crediting any paper that names an accuracy confound, would move it, which is why the rule is stated explicitly rather than the count asserted alone.

The "analysed as a selection bias?" coding rule. A paper counts as analysing the effect only if it both (i) names the correctness-conditioned comparison as a possible source of distortion in the faithfulness *estimate* and (ii) takes an analytical step with respect to it, a correction, a sensitivity analysis, or an explicit statement that the reported number is conditional and how the conditioning could move it. A passing remark that an accuracy confound exists, without either step, is an incidental mention and does not count. The strict reading is deliberate: our claim is that no paper treats this as a selection problem, so the coding must not credit a paper for noticing the confound while leaving the estimate unqualified.

The distinction is not hypothetical. Bentham et al. [2024] states plainly that correctness filtering is central and "can inflate or alter interpretation of faithfulness estimates" (the closest any audited paper comes, and why we credit it in Section 2 with the accuracy confound), yet it offers no correction, no sensitivity analysis, and no unqualified reporting of the filtered number; it codes as an incidental mention. Appendix H.1 records the rule's application and the authors' audit notes retain the evidence behind each judgement. Those working notes are not part of the public release. Most in-scope notes are catalog-derived paraphrases rather than verbatim page-and-section quotations, and the rows were not independently double-coded, so the audit is a transparent single-author coding result rather than an independently registered systematic review; these limits are stated explicitly.

5 EXPERIMENTAL DESIGN

We design a controlled framework that (i) provides ground-truth causal structure, (ii) embeds configurable shortcut channels, and (iii) compares the correctness-filtered estimator F_{sel} against the unfiltered F_{pop}.

5.1 Procedural Task Design

Intervention-based faithfulness evaluation requires knowing the correct counterfactual answer after a step is edited. Natural-language tasks generally do not provide a unique, verifiable downstream answer after such an edit. We therefore construct procedural tasks where a symbolic executor replays the dependency graph and produces the oracle counterfactual deterministically [Lyu et al., 2023, Hase and Potts, 2026]; we use this controlled setting to study *measurement bias*, not to claim that the generators represent natural-language reasoning.

We implement three structurally distinct families, each with a dependency graph, editable intermediate values, deterministic oracles, and semantics-preserving placebo edits: **arithmetic DAGs** of integer operations (depth 2–4); **Boolean circuits** of logical gates (depth 2–8); and **graph reachability** on 5–12 nodes. Each generator produces 100 candidate instances per family with deterministic seeds; the headline experiments use a subset (Table 3). A faithful baseline ($n = 50$) validates the apparatus, giving $F_{\text{pop}} = F_{\text{sel}} = 1.000$ (Appendix G).

Planted shortcut channels. Purely procedural tasks are structurally shortcut-proof: every step depends on prior computed values, so $S \approx 0$ and $\Delta = 0$. To create a reachable shortcut regime, we augment instances with a planted cue whose validity is controlled by q, the probability that the cue agrees with the oracle answer. The cue is not updated under intervention, so a cue-reliant model can pass factual correctness while failing counterfactual compliance. The generator control q is distinct from the theoretical $p_s = P(S{=}1 \mid D)$, which denotes the prevalence of shortcut-capable behavior in the SCM and is set directly in the implementation check.

5.2 All-Trace Intervention Protocol

We strictly separate two roles: the language model under test generates traces and answers; the symbolic executor (oracle) computes ground-truth values by replaying the dependency graph. For each instance i the model processes three conditions independently: a **factual condition** on the original task, recording correctness $C_i = \mathbb{1}[\hat{y}_i^{\text{fact}} = y_i^*]$; a **causal intervention condition** in which one intermediate graph value is changed and the oracle computes the correct counterfactual $y_i^{\text{cf}*}$, recording compliance $A_i = \mathbb{1}[\hat{y}_i^{\text{cf}} = y_i^{\text{cf}*}]$; and a **placebo condition** applying a semantics-preserving edit, recording $A_i^{\text{plac}} = \mathbb{1}[\hat{y}_i^{\text{plac}} = y_i^*]$.

Each condition is presented as a standalone task. Non-interventionable traces (missing or unparseable answers) are scored non-compliant ($A_i = 0$) in our headline convention; Appendix P reports a null-as-missing sensitivity analysis and documents that this convention inflates rather than deflates the estimated gap. The placebo condition was used to construct and validate the instance set and is not reported per-model.

5.3 Evaluation Metrics

The primary estimands are:

$$F_{\text{pop}} = \frac{1}{n}\sum_{i=1}^{n} A_i, \qquad F_{\text{sel}} = \frac{\sum_{i=1}^{n} A_i \cdot C_i}{\sum_{i=1}^{n} C_i}, \qquad \Delta = F_{\text{sel}} - F_{\text{pop}}. \tag{20}$$

All confidence intervals use the nonparametric bootstrap ($B = 10{,}000$ resamples of *instances*, the unit at which the model was queried). The pre-specified primary analysis is the pooled Cochran–Mantel–Haenszel test (a single hypothesis, so no multiplicity adjustment is required); individual model p-values are descriptive. All model runs for the headline results use temperature $= 0$ (greedy decoding).

6 Applying the Framework to Real Language Models

We now apply the framework of Section 3 to real models, asking whether the correctness filter biases measured faithfulness and which pathway carries the bias. We run the all-trace protocol on four sub-ceiling models from two providers (ZhipuAI, DeepSeek), plus two 100%-accuracy boundary controls from two further providers (Alibaba, Google), on arithmetic DAG tasks (Table 1).

Table 1: Empirical selection bias on real language models. Four sub-ceiling models from two providers; two boundary controls from two further providers. Results describe specific serving deployments, not model families; deployment-level variation is documented in Appendix R. All counts use the null-as-non-compliant convention (Section 5.2), so each row's denominator is the full n; the per-model (A, C) cells are printed in Table 7. The controls have $P(C{=}1) = 1$ by construction and $\Delta = 0$ follows algebraically (Appendix O). $^*p < 0.05$.

Model	n	Acc.	F_{pop}	F_{sel}	Δ (pp)	95% CI (pp)	p
glm-5.2	100	79.0%	0.440	0.456	+1.6	[−3.4, +6.6]	0.540
deepseek-v4-pro	200	67.0%	0.555	0.579	+2.4	[−2.6, +7.2]	0.337
glm-5.3	**1000**	**51.5%**	**0.507**	**0.581**	**+7.4**	**[+4.3, +10.4]**	**<0.0001***
deepseek-v4-flash	**300**	**48.3%**	**0.527**	**0.586**	**+6.0**	**[+0.2, +11.8]**	**0.046***
Boundary controls ($P(C{=}1) = 1$, filter vacuous):							
qwen3.8-max	20	100%	0.600	0.600	0.0	N/A	N/A
Gemini 3.5 Flash Lite	17	100%	0.529	0.529	0.0	N/A	N/A

The filter biases measured faithfulness. GLM-5.3 shows a significant gap of +7.4pp ($p < 0.0001$, $n = 1000$; Figure 5); deepseek-v4-flash is marginal (+6.0pp, $p = 0.046$, $n = 300$).[1] The remaining two models show directionally positive but non-significant gaps; a Cochran–Mantel–Haenszel test pooled across the four gives OR $= 1.70$ (Appendix L). Leave-one-out analysis shows that the pooled effect survives removing any single model: excluding GLM-5.3, the largest contributor, reduces the pooled z from 5.19 to 2.24 ($p = 0.025$), while still showing GLM-5.3's influence. The two 100%-accuracy controls show exactly zero gap because $P(C{=}1) = 1$ makes the filter vacuous; they provide an algebraic boundary check and add no empirical evidence. As a sensitivity check, excluding the seed-3000

[1] Individual model p-values are descriptive and two-sided (Pearson χ^2, $df{=}1$); the pre-specified primary analysis is the pooled CMH test.

batch entirely (the only one requiring deduplication, contributing 100 of the 1000 rows), the remaining $n = 900$ rows (seeds 6000, 12000, 13000; zero deduplication) give a gap of $+7.26$pp ($\chi^2 = 20.54$, $p < 0.0001$), confirming that the deduplicated batch does not drive the result. **The variation associated with the measured proxy.** Conditioning on $C = 1$ selects instances with fewer interventionable operations, where compliance is higher. We use two summaries with different interpretations. First, the standard covariance decomposition equation 18, stratified on the interventionable-operation count, assigns 42.8% (GLM-5.3) and 21.9% (flash) of the gap to between-stratum variation *in that proxy*; the remainder is within-stratum and *unattributed* (Table 8). The GLM-5.3 share is the precisely-estimated case ($n = 1000$); flash's $n = 300$ strata are small enough that its share is imprecise, with a bootstrap interval spanning zero. The share is also covariate-dependent: oracle-displacement magnitude gives 26.6% on the same GLM-5.3 records, so 42.8% is the share captured by operation count, not a causal contribution of latent difficulty. Second, inverse-propensity weighting on the same covariate reweights selected rows toward the observed population, removing 23–34% of the raw gap across two propensity specifications. Both summaries leave the majority unexplained. Shortcut reliance is indistinguishable from zero here: the trace-matching detector fires on 0% of instances under neutral framing, and the adversarial elicitation raises that to 10% in GLM-5.3 only. The sharper cue-correctness contrast is reported on complete cases ($n = 291$), yielding -1.6pp ($z = -0.24$, $p = 0.81$); it uses 112/208 correct-cue rows and 46/83 incorrect-cue rows, excluding 9 null rows. The shortcut branch is empty for these deployments under neutral framing, and the prospective analysis of Section 3.4 is not tested on the headline runs. A targeted follow-up reaches the branch on several deployments under an explicit-reliance framing (Appendix M); those are separate regimes, not pooled with these gaps. These gaps describe the specific serving deployments that answered our requests, not the model families (Appendix R).

Which estimand the headline number is. Table 1 reports the *marginal* $\Delta = F_{\text{sel}} - F_{\text{pop}}$; the theory identifies the stratum-specific $\Delta^{(d)}$ (Appendix D), and the marginal is a $P(C{=}1 \mid D{=}d)$-weighted mixture of those. They differ materially on GLM-5.3: the marginal is $+7.4$pp, while selected-weighted difficulty standardization, $\Delta_{\text{std},C} \coloneqq \sum_d P(D{=}d \mid C{=}1)\,(F^{(d)}_{\text{sel}} - F^{(d)}_{\text{pop}})$, gives $+4.2$pp, not the naive average ($+4.81$pp). The 3.2pp difference is the between-stratum term of equation 18, not a within-stratum effect.

The closed-form mixture is not the estimand for these runs. The mixture form of §3.4 was derived for a population with a live shortcut branch. Our headline runs have none (measured shortcut use 0%), so the form does not predict their gaps. Solving each model's own F_{pop} and $P(C{=}1)$ for an implied p_s under Corollary 1 gives $+0.625$, $+0.258$, $+0.016$ and -0.092 for glm-5.2, v4-pro, GLM-5.3 and flash: all near zero, one negative, hence not a valid probability. Plugging those back gives predicted gaps of 11.7–56.3pp against observed 1.6–7.4pp, a factor of 6.4–11.4 (Table 2). This is a boundary of the special case, not a failure of the observed (A, C) estimand. The empirical result, that correctness conditioning biases the estimate and that operation-count strata capture a measurable component, does not depend on the mixture form.

7 Conclusion

Correctness conditioning is a selection step in CoT faithfulness evaluation, and the identity $\Delta = \text{Cov}(A, C)/P(C{=}1)$ quantifies the resulting gap without assigning it to a mechanism. A structured audit found 16 of 18 published evaluations report a conditioning-dependent

Table 2: Over-prediction of the closed-form mixture, with its inputs printed. Each model's implied p_s is obtained from its own F_{pop} and accuracy $P(C{=}1)$ under Corollary 1; the predicted gap then follows from the same expression. The flash row's implied p_s is negative, so its predicted value is not a valid probability and is shown for completeness only. Ratios span 6.4–11.4×.

Model	F_{pop} ($= p_r$)	Acc. ($= P(C{=}1)$)	Implied p_s	Predicted Δ (pp)	Observed Δ (pp)	Ratio
glm-5.2	0.440	79.0%	$+0.625$	11.7	1.6	$7.3\times$
deepseek-v4-pro	0.555	67.0%	$+0.258$	27.3	2.4	$11.4\times$
GLM-5.3	0.507	51.5%	$+0.016$	47.7	7.4	$6.4\times$
deepseek-v4-flash	0.527	48.3%	-0.092	56.3	6.0	$9.4\times$

metric. On four sub-ceiling deployments, the primary run shows a $+7.4$pp gap (GLM-5.3, $n = 1000$, $p < 0.0001$); operation-count strata account for 42.8% of that gap in this dataset, while the remainder is unattributed. The evidence comes from one synthetic task family, only two of four sub-ceiling models have individually significant gaps, and the collider pathway is inactive under neutral framing. When counterfactual compliance is observable for every instance, all-trace evaluation directly estimates F_{pop}.

Reproducibility Statement

The procedural task generators use deterministic seed streams, and the oracle computations are symbolic with exact outputs. Simulation parameters, estimands, prompts, and scoring conventions are specified in the main text and appendices. Headline model calls use temperature $= 0$ (greedy decoding); when a serving endpoint returned non-identical responses at that setting, we retained the first completed response for the reported run. The headline experiments were conducted in August 2026, and the exploratory framing probe was conducted in September 2026. The reported sample sizes are the final counts after removing repeated queries within a model–seed–instance key; temporary seed labels and intermediate batch exports are implementation bookkeeping and are intentionally omitted.

The public repository reports the aggregate results in the paper tables and figures and contains the website source, paper figures, paper PDF, and deterministic exhibit checks. The public release does not claim to provide a row-level rerun of the deployment-bound experiments: provider request metadata, raw model responses, temporary logs, internal audit ledgers, and credential-bound generators are excluded. This boundary keeps the release readable and prevents intermediate experiment labels from being mistaken for scientific variables.

Data and Code Availability

The website source, paper figures, PDF, and release checks are available in the CCS-CoTFE GitHub repository at `https://github.com/CCS-CoTFE/CCS-CoTFE`, with the interactive exhibit published at `https://ccs-cotfe.github.io/CCS-CoTFE/`. The paper source bundle contains the LaTeX files and the figure assets needed to compile the manuscript. The repository also contains the project film and the public exhibit code.

No supplementary-material submission or Hugging Face dataset is claimed for this

release. Raw API logs, model completions, provider request identifiers, intermediate experiment exports, and the working literature-audit ledger are intentionally excluded from both the public repository and the paper source bundle. Headline values can be inspected in the printed tables and figures; exact row-level reanalysis would require a separately released, anonymized data snapshot.

Ethics Statement

This work studies the statistical properties of evaluation methodologies and does not involve human subjects, private data, or model training. All language model API calls use publicly available commercial endpoints under standard terms of service. The procedural tasks are synthetically generated and contain no sensitive content.

References

Iván Arcuschin, Jett Janiak, Robert Krzyzanowski, Senthooran Rajamanoharan, Neel Nanda, and Arthur Conmy. Chain-of-thought reasoning in the wild is not always faithful. *arXiv preprint arXiv:2503.08679*, 2025.

Pepa Atanasova, Oana-Maria Camburu, Christina Lioma, Jakob Grue Simonsen, Thomas Lukasiewicz, and Isabelle Augenstein. Faithfulness tests for natural language explanations. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*, pages 283–294, 2023.

Bowen Baker, Joost Huizinga, Leo Gao, Zehao Dou, Melody Y. Guan, Aleksander Madry, Wojciech Zaremba, Jakub Pachocki, and David Farhi. Monitoring reasoning models for misbehavior and the risks of promoting obfuscation. *arXiv preprint arXiv:2503.11926*, 2025.

Sriram Balasubramanian, Samyadeep Basu, and Soheil Feizi. A closer look at bias and chain-of-thought faithfulness of large (vision) language models. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 13406–13439, 2025.

Elias Bareinboim and Judea Pearl. Controlling selection bias in causal inference. In *Proceedings of the Fifteenth International Conference on Artificial Intelligence and Statistics*, volume 22 of *Proceedings of Machine Learning Research*, pages 100–108, 2012.

Elias Bareinboim, Jin Tian, and Judea Pearl. Recovering from selection bias in causal and statistical inference. In *Proceedings of the Twenty-Eighth AAAI Conference on Artificial Intelligence*, pages 2410–2416, 2014.

Oliver Bentham, Nathan Stringham, and Ana Marasović. Chain-of-thought unfaithfulness as disguised accuracy. *Transactions on Machine Learning Research*, 2024.

Joseph Berkson. Limitations of the application of fourfold table analysis to hospital data. *Biometrics Bulletin*, 2(3):47–53, 1946.

Paul C. Bogdan, Uzay Macar, Neel Nanda, and Arthur Conmy. Thought anchors: Which LLM reasoning steps matter? *arXiv preprint arXiv:2506.19143*, 2025.

Yanda Chen, Joe Benton, Ansh Radhakrishnan, Jonathan Uesato, Carson Denison, John Schulman, Arushi Somani, Peter Hase, Misha Wagner, Fabien Roger, Vlad Mikulik, Samuel R. Bowman, Jan Leike, Jared Kaplan, and Ethan Perez. Reasoning models don't always say what they think. *arXiv preprint arXiv:2505.05410*, 2025.

Corinna Cortes, Mehryar Mohri, Michael Riley, and Afshin Rostamizadeh. Sample selection bias correction theory. In *Algorithmic Learning Theory, 19th International Conference (ALT)*, volume 5254 of *Lecture Notes in Computer Science*, pages 38–53, 2008.

Ran El-Yaniv and Yair Wiener. On the foundations of noise-free selective classification. *Journal of Machine Learning Research*, 11(53):1605–1641, 2010.

Felix Elwert and Christopher Winship. Endogenous selection bias: The problem of conditioning on a collider variable. *Annual Review of Sociology*, 40:31–53, 2014.

Jiarun Fu, Lizhong Ding, Hao Li, Pengqi Li, Qiuning Wei, and Xu Chen. Unveiling and causalizing CoT: A causal perspective. *arXiv preprint arXiv:2502.18239*, 2025.

Yonatan Geifman and Ran El-Yaniv. Selective classification for deep neural networks. In *Advances in Neural Information Processing Systems*, volume 30, pages 4878–4887, 2017.

Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A. Wichmann. Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2(11):665–673, 2020.

Yoav Gur-Arieh, Ana Marasović, and Mor Geva. Faithfulness metrics don't measure faithfulness: A meta-evaluation with ground truth. *arXiv preprint arXiv:2605.25052*, 2026.

Peter Hase and Christopher Potts. Counterfactual simulation training for chain-of-thought faithfulness. In *Conference on Language Modeling (COLM)*, 2026.

James J. Heckman. Sample selection bias as a specification error. *Econometrica*, 47(1): 153–161, 1979.

D. G. Horvitz and D. J. Thompson. A generalization of sampling without replacement from a finite universe. *Journal of the American Statistical Association*, 47(260):663–685, 1952.

Jinghan Jia, Joe Benton, and Eric Easley. Faithfulness as information flow: Evaluating and training faithful chain-of-thought reasoning. *arXiv preprint arXiv:2605.24286*, 2026.

Zhijing Jin, Yuen Chen, Felix Leeb, Luigi Gresele, Ojasv Kamal, Zhiheng Lyu, Kevin Blin, Fernando Gonzalez Adauto, Max Kleiman-Weiner, Mrinmaya Sachan, and Bernhard Schölkopf. CLadder: Assessing causal reasoning in language models. In *Advances in Neural Information Processing Systems*, volume 36, pages 31038–31065, 2023.

Himabindu Lakkaraju, Jon Kleinberg, Jure Leskovec, Jens Ludwig, and Sendhil Mullainathan. The selective labels problem: Evaluating algorithmic predictions in the presence of unobservables. In *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 275–284, 2017.

Tamera Lanham, Anna Chen, Ansh Radhakrishnan, Benoit Steiner, Carson Denison, Danny Hernandez, Dustin Li, Esin Durmus, Evan Hubinger, Jackson Kernion, Kamilė Lukošiūtė, Karina Nguyen, Newton Cheng, Nicholas Joseph, Nicholas Schiefer, Oliver Rausch, Robin Larson, Sam McCandlish, Sandipan Kundu, Saurav Kadavath, Shannon Yang, Thomas Henighan, Timothy Maxwell, Timothy Telleen-Lawton, Tristan Hume, Zac Hatfield-Dodds, Jared Kaplan, Jan Brauner, Samuel R. Bowman, and Ethan Perez. Measuring faithfulness in chain-of-thought reasoning. *arXiv preprint arXiv:2307.13702*, 2023.

Wenkai Li, Fan Yang, Ananya Hazarika, Shaunak A. Mehta, and Koichi Onoue. When reasoning traces become performative: Step-level evidence that chain-of-thought is an imperfect oversight channel. *arXiv preprint arXiv:2605.11746*, 2026.

Qing Lyu, Shreya Havaldar, Adam Stein, Li Zhang, Delip Rao, Eric Wong, Marianna Apidianaki, and Chris Callison-Burch. Faithful chain-of-thought reasoning. In *International Joint Conference on Natural Language Processing and Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics (IJCNLP-AACL)*, 2023.

Emily McMilin. Selection collider bias in large language models. In *UAI 2022 Workshop on Causal Representation Learning (CRL)*, 2022. arXiv:2208.10063.

Avni Mittal and Rauno Arike. C^2-Faith: Benchmarking LLM judges for causal and coverage faithfulness in chain-of-thought reasoning. In *Proceedings of the Fifth Workshop on Generation, Evaluation and Metrics (GEM)*, 2026.

Karthika Mohan and Judea Pearl. Graphical models for processing missing data. *Journal of the American Statistical Association*, 116(534):1023–1037, 2021.

Letitia Parcalabescu and Anette Frank. On measuring faithfulness or self-consistency of natural language explanations. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, pages 6048–6089, 2024.

Manas Pathak, Xingyao Chen, Shuozhe Li, Amy Zhang, and Liu Leqi. Filtered reasoning score: Evaluating reasoning quality on a model's most-confident traces. In *Conference on Language Modeling (COLM)*, 2026.

Debjit Paul, Robert West, Antoine Bosselut, and Boi Faltings. Making reasoning matter: Measuring and improving faithfulness of chain-of-thought reasoning. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 15012–15032, 2024.

Judea Pearl. *Causality: Models, Reasoning, and Inference*. Cambridge University Press, 2nd edition, 2009.

Ansh Radhakrishnan, Karina Nguyen, Anna Chen, Carol Chen, Carson Denison, Danny Hernandez, Esin Durmus, Evan Hubinger, Jackson Kernion, Kamilė Lukošiūtė, Newton Cheng, Nicholas Joseph, Nicholas Schiefer, Oliver Rausch, Sam McCandlish, Sheer El Showk, Tamera Lanham, Tim Maxwell, Venkatesa Chandrasekaran, Zac Hatfield-Dodds, Jared Kaplan, Jan Brauner, Samuel R. Bowman, and Ethan Perez. Question decomposition improves the faithfulness of model-generated reasoning. *arXiv preprint arXiv:2307.11768*, 2023.

Paul R. Rosenbaum. *Observational Studies*. Springer, 2nd edition, 2002.

Donald B. Rubin. Inference and missing data. *Biometrika*, 63(3):581–592, 1976.

Xu Shen, Song Wang, Zhen Tan, Laura Yao, Xinyu Zhao, Kaidi Xu, Xin Wang, and Tianlong Chen. FaithCoT-bench: Benchmarking instance-level faithfulness of chain-of-thought reasoning. In *International Conference on Learning Representations*, 2026.

Lin Shi, Chiyu Ma, Wenhua Liang, Xingjian Diao, Weicheng Ma, and Soroush Vosoughi. Judging the judges: A systematic study of position bias in LLM-as-a-judge. In *Proceedings of the 14th International Joint Conference on Natural Language Processing and the 4th Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics*, pages 292–314, 2025. doi: 10.18653/v1/2025.ijcnlp-long.18.

Noah Y. Siegel, Oana-Maria Camburu, Nicolas Heess, and Maria Perez-Ortiz. The probabilities also matter: A more faithful metric for faithfulness of free-text explanations in large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 530–546, 2024.

Zayne Sprague, Fangcong Yin, Juan Diego Rodriguez, Dongwei Jiang, Manya Wadhwa, Prasann Singhal, Xinyu Zhao, Xi Ye, Kyle Mahowald, and Greg Durrett. To CoT or not to CoT? chain-of-thought helps mainly on math and symbolic reasoning. *arXiv preprint arXiv:2409.12183*, 2024.

Yuxi Sun, Aoqi Zuo, Haotian Xie, Wei Gao, Mingming Gong, and Jing Ma. FACT-E: Causality-inspired evaluation for trustworthy chain-of-thought reasoning. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 20275–20293, 2026.

Sree Harsha Tanneru, Dan Ley, Chirag Agarwal, and Himabindu Lakkaraju. On the hardness of faithful chain-of-thought reasoning in large language models. *arXiv preprint arXiv:2406.10625*, 2024.

Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. Language models don't always say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in Neural Information Processing Systems*, volume 36, 2023.

Martin Tutek, Fateme Hashemi Chaleshtori, Ana Marasović, and Yonatan Belinkov. Measuring chain of thought faithfulness by unlearning reasoning steps. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 9935–9960, 2025.

Jesse Vig, Sebastian Gehrmann, Yonatan Belinkov, Sharon Qian, Daniel Nevo, Yaron Singer, and Stuart Shieber. Investigating gender bias in language models using causal mediation analysis. In *Advances in Neural Information Processing Systems*, volume 33, pages 12388–12401, 2020.

Keheng Wang, Feiyu Duan, Sirui Wang, Peiguang Li, Yunsen Xian, Chuantao Yin, Wenge Rong, and Zhang Xiong. Knowledge-driven CoT: Exploring faithful reasoning in LLMs for knowledge-intensive question answering. *arXiv preprint arXiv:2308.13259*, 2023.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc V. Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems*, volume 35, 2022.

Richard J. Young. Measuring faithfulness depends on how you measure: Classifier sensitivity in LLM chain-of-thought evaluation. *arXiv preprint arXiv:2603.20172*, 2026.

Bianca Zadrozny. Learning and evaluating classifiers under sample selection bias. In *Proceedings of the Twenty-First International Conference on Machine Learning (ICML)*, pages 903–910, 2004.

Kerem Zaman and Shashank Srivastava. A causal lens for evaluating faithfulness metrics. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 29425–29449, 2025.

A Experimental Configuration

Empirical instance counts are reported per batch because task families and experimental phases use different batch sizes. The implementation check uses small deterministic batches across three synthetic task families. The real-model study uses arithmetic DAGs with a primary run of $n = 1000$ and additional sub-ceiling deployments of $n = 300$, $n = 200$, and $n = 100$, together with two 100%-accuracy boundary controls. Repeated queries within a model–seed–instance key are removed before aggregation. Exact seed labels and temporary batch identifiers are implementation bookkeeping rather than scientific variables, so they are not part of the paper's public description.

The arithmetic DAG generator uses a depth range of $[2, 4]$, a value range of $[1, 20]$, and a planted-cue validity control of $q = 0.7$. This generator control is distinct from the theoretical shortcut prevalence p_s used in the implementation check.

The prompt template for the real-model runs is:

```
Solve this step by step.  Show your work.

{task prompt}

Your final answer (just the number):
```

The counterfactual condition replaces the intervened step line with the new value (e.g. `S3 = 12 (given)`) while keeping the rest of the task unchanged; no system-level objective or persona prompt was used in the empirical runs. Persona-modulated agents in simulation (Section B) are deterministic compliance models: for each instance, the agent answers correctly with probability π_{rs} and complies with intervention with probability η_{rs}, drawn from the closed-form cell probabilities. They are mathematical validation devices and do not correspond to real language models.

B Implementation Consistency Check

This section checks our implementation against our derivations; it is not evidence for the theory. The agents are drawn from the same closed-form cell probabilities the propositions derive from, so agreement is guaranteed by construction and the assumptions are imposed rather than tested. Sweeping $(p_r, p_s) \in [0, 1]^2$ gives a mean absolute error of 0.006 against Corollary 1, reproducing both analytic predictions: the gap decreases monotonically in p_s (Corollary 2) and is hump-shaped in p_r at the analytic location $p_r^* \approx 0.456$ (Figure 2, Appendix C).

The mechanism is not tied to one simulated task family. The three procedural families have different structures (integer arithmetic, Boolean logic, and graph reachability) and different reliance levels. Table 3 reports the gap for each at a common theoretical shortcut prevalence $p_s = 0.7$. The gaps agree to within 0.6pp (mean $+8.2$pp, band $+7.9$ to $+8.5$pp) even though the per-family reliance p_r ranges from 0.33 to 0.58. This cross-family implementation check does not expand the single-family real-model evidence; it shows only that the closed-form mechanism is not specific to arithmetic-DAG syntax.

Family	n	p_r	F_{pop}	F_{sel}	Gap (pp)
Arithmetic DAG	30	0.33	0.333	0.417	+8.3
Boolean circuit	50	0.58	0.580	0.659	+7.9
Graph reachability	50	0.52	0.520	0.605	+8.5
Mean					**+8.2**

Table 3: Simulation gap across three structurally distinct task families ($p_s = 0.7$ throughout; here p_s is the theoretical shortcut prevalence, and p_r varies by family and equals F_{pop} under strict reliance). The gaps agree to within 0.6pp despite a 0.25-wide spread in p_r. Full per-family detail in Appendix C.

Under hard-task parameters, Model $\mathcal{A}$ ($p_r = 0.45$, $p_s = 0.90$) has $F_{\text{pop}} = 0.450$ but $F_{\text{sel}} = 0.504$, while Model $\mathcal{B}$ ($p_r = 0.70$, $p_s = 0.01$) has $F_{\text{pop}} = 0.700$ but $F_{\text{sel}} = 0.308$: conditioning reverses the ranking, with 40.6% of candidates reversing at the extreme setting (Appendix K). These are the conditions Section 3.4 flags as prospective: they hold at $S > 0$, which our real-model runs do not reach.

C Simulation: Cross-Family Detail

Table 3 (Section B) reports the per-family simulation results. The theoretical shortcut prevalence is $p_s = 0.7$ for every row; this is distinct from the planted-cue validity control q in Section 5.1. The per-family CoT-reliance level p_r differs and equals F_{pop} under strict reliance. Boolean circuit and graph tasks use $n = 50$ for richer structural variation; arithmetic DAGs use $n = 30$ due to simpler compositional structure. The three families share the intervention protocol and oracle design of Section 5; what differs is the dependency structure, which is the point of the comparison.

D Estimand: Within-Stratum vs. Marginal

The paper uses two estimands and it matters which result speaks to which. The theory's Propositions 1, 3 and 4 are stated at fixed difficulty $D = d$ and identify the *stratum-specific* $\Delta^{(d)}$. Conditioning on d is what makes $R \perp S \mid D$ coherent and what permits the T-marginalization; it also, by construction, removes the difficulty-confounding pathway from $\Delta^{(d)}$, since D no longer varies. Proposition 2 is the exception: it is stated *marginally* in D, and is the bridge to the real-model runs, where R and S are latent and cannot be conditioned on.

Table 1 instead reports the *marginal* $\Delta = F_{\text{sel}} - F_{\text{pop}}$, which is a $P(C{=}1 \mid D{=}d)$-weighted mixture of the $\Delta^{(d)}$. That mixture is *not* identified by the within-stratum result:

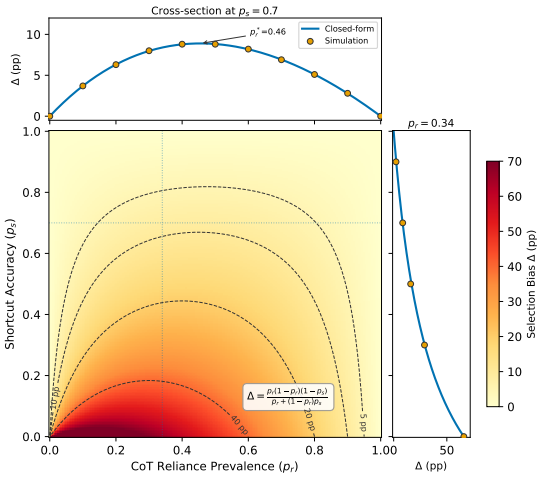


Figure 2: Selection bias surface $\Delta(p_r, p_s)$ from the closed-form (heatmap) with simulation points (orange) in the marginals. Hump-shaped in p_r (top), monotonically decreasing in p_s (right). MAE $= 0.006$.

difficulty enters it directly, which is why we read the real-model gaps as consistent with difficulty confounding rather than as an instance of the mixture form.

The practical consequence is stated in the main text: the within-stratum analysis (Appendix J) decomposes the marginal gap by operation count, and its finding that gaps persist within strata is what tells us the marginal Δ is not purely between-stratum. Neither route tests the collider pathway, because the shortcut branch is empty in these runs.

Magnitude reconciliation. Two magnitudes look contradictory unless the shortcut branch is made explicit, and both discrepancies have one source. First, setting $p_s = 0$ (matching our measured 0% shortcut use) makes Corollary 1 degenerate to $\Delta = 1 - P(C{=}1)$, monotone with no interior maximum, predicting $\sim$48pp at GLM-5.3's 52% accuracy against $+7.4$pp observed. Second, the implementation check sets the theoretical mixture parameter to $p_s = 0.7$, which implies an accuracy floor of 0.70; the measured 48–52% deployments are outside that special-case population. In both cases the mixture form is applied to a population with a live shortcut branch while the empirical runs have none.

Quantifying it directly: solving each model's own F_{pop} and accuracy for an implied p_s gives $+0.625$, $+0.258$, $+0.016$, and -0.092 (glm-5.2, v4-pro, GLM-5.3, flash)—all near zero and one negative, which is not a valid probability and independently corroborates the measured 0% shortcut use. Table 2 (Section 6) prints the inputs to that calculation so the predicted/observed ratios can be re-derived directly; Section 6 states the resulting tension with the theory.

The mixture special cases are therefore scoped to the simulation and to the idealised design analysis; the real-model gaps are reported as a separate empirical finding about the operation-count covariate.

Two consequences follow. The accuracy-axis peak sits at $P^* = \sqrt{p_s}$, so with the implementation-check value $p_s = 0.7$ it lies at 83.7%: the four sub-ceiling deployments (48–79%) lie below it, while the two ceiling controls (100%) lie above. No tested deployment sits near the peak. The ceiling controls cannot localise it because $P(C{=}1) = 1$ makes $\Delta = 0$ an algebraic identity. Probing a peak near 50% accuracy would require $p_s \approx 0.25$, which is

a different theoretical regime from the one used in the implementation check. The accuracy band in Figure 5 is therefore descriptive, not a prediction.

E Proof Details

E.1 Corollary 4: Hump on the Accuracy Axis

With $r \coloneqq \sqrt{p_s}$ and strict reliance, Corollary 3 locates the hump at $p_r^* = r/(1+r)$ in CoT-reliance space. Substituting into Corollary 1: $p_r^*(1-p_r^*) = r/(1+r)^2$, and $p_r^* + (1-p_r^*)p_s = r/(1+r) + r^2(1+r)^{-1} = r$. Hence

$$\Delta(P^*) = \frac{r(1-p_s)}{(1+r)^2\, r} = \frac{1-r^2}{(1+r)^2} = \frac{1-r}{1+r}, \qquad P(p_r^*) = p_r^* + (1-p_r^*)p_s = r = \sqrt{p_s}. \tag{21}$$

Since P is affine and strictly increasing in p_r, the unique interior maximum in p_r maps to the unique interior maximum in P. Attainable accuracy spans $[p_s, 1]$; the peak $P^* = \sqrt{p_s}$ lies in that interval for all $p_s \in (0, 1)$.

Proofs of Propositions 1 and 2 are given in the main text. The collider pathway derivations (Propositions 3 and 4, Corollaries 1–4) are in Appendix F. This appendix contains additional derivation details.

F Collider Pathway: Full Derivations

The following derivations support the prospective analysis summarized in Section 3.4. They require an active shortcut branch ($S > 0$), which none of our headline deployments exhibits under neutral framing.

The cleanest special case sets $\eta_{00} = \eta_{01} = 0$ and $\eta_{10} = \eta_{11} = 1$: the model's answer changes under intervention if and only if it genuinely relies on the CoT. Then $F_{\text{pop}} = p_r$. Define the stratum-specific correctness rates

$$g_1(p_s) \coloneqq P(C{=}1 \mid R{=}1) = (1{-}p_s)\,\pi_{10} + p_s\,\pi_{11}, \tag{22}$$

$$g_0(p_s) \coloneqq P(C{=}1 \mid R{=}0) = (1{-}p_s)\,\pi_{00} + p_s\,\pi_{01}. \tag{23}$$

Proposition 3 (Strict-Reliance Gap)**.** *Under the strict-reliance assumption, $F_{\text{sel}} = P(R{=}1 \mid C{=}1)$, and the gap is:*

$$\Delta = \frac{p_r(1-p_r)(g_1 - g_0)}{p_r\, g_1 + (1-p_r)\, g_0}. \tag{24}$$

The sign of Δ is determined by $g_1 - g_0$: correctness conditioning overestimates faithfulness when CoT reliance makes correctness more likely ($g_1 > g_0$), and underestimates it otherwise.

Proof. Under strict reliance, $F_{\text{sel}} = P(R{=}1 \mid C{=}1) = p_r g_1 / [p_r g_1 + (1{-}p_r) g_0]$ by Bayes' rule. Subtracting $F_{\text{pop}} = p_r$ and simplifying yields equation 24. □

Corollary 1 (Simplified Gap Under Perfect Faithful Reasoning)**.** *When faithful reasoning is always correct ($\pi_{10} = \pi_{11} = 1$), shortcut-only correctness is certain ($\pi_{01} = 1$), and the neither-route cell is always wrong ($\pi_{00} = 0$), we have $g_1 = 1$ and $g_0 = p_s$, giving:*

$$\Delta = \frac{p_r(1-p_r)(1-p_s)}{p_r + (1-p_r)\, p_s}. \tag{25}$$

This simplified form is a theoretical mixture for an active shortcut branch. The planted-cue generator in Section 5.1 uses the separate validity control q and does not estimate p_s in the real-model runs. For multiple-choice tasks with k options, $\pi_{00} = 1/k > 0$ and the general formula equation 24 applies; the simplified form remains a close approximation when $1/k \ll p_s$.

Corollary 2 (Monotone Decrease in Shortcut Prevalence)**.** *For fixed $p_r \in (0,1)$, $\partial\Delta/\partial p_s < 0$: the gap equation 25 is strictly decreasing in shortcut prevalence.*

Proof. Write $\Delta(p_s) = a\,b\,(1-p_s)/(a+b\,p_s)$ with $a = p_r > 0$ and $b = 1-p_r > 0$. Differentiating: $\partial\Delta/\partial p_s = -a\,b\,(a+b)/(a+b\,p_s)^2 < 0$. □

Corollary 3 (Hump-Shaped Prevalence Response)**.** *For fixed $p_s \in (0,1)$, $\Delta(p_r)$ from equation 25 is zero at $p_r = 0$ and $p_r = 1$, with a unique interior maximum at $p_r^* = \sqrt{p_s}/(1+\sqrt{p_s})$.*

Proof. Setting $d\Delta/dp_r = 0$: the first-order condition yields $p_r^2 = p_s(1-p_r)^2$, hence $p_r/(1-p_r) = \sqrt{p_s}$, giving $p_r^* = \sqrt{p_s}/(1+\sqrt{p_s})$. The boundary values $\Delta(0) = \Delta(1) = 0$ confirm this is a maximum. □

Corollary 4 (Hump on the Accuracy Axis)**.** *Under the strict-reliance assumptions of Corollary 1, with $P \coloneqq P(C{=}1)$, the bias has a unique interior maximum in accuracy at $P^* = \sqrt{p_s}$, where $\Delta(P^*) = (1-\sqrt{p_s})/(1+\sqrt{p_s})$ (proof in Appendix E). Accuracy below P^* lies on the rising limb, above it on the falling limb.*

Proposition 4 (Ranking Reversal)**.** *Consider two models $\mathcal{A}$ and $\mathcal{B}$ in the strict-reliance case with shared correctness parameters (π_{rs}). If $p_{r,A} < p_{r,B}$ (model $\mathcal{A}$ is less faithful at the population level), a ranking reversal $F_{\text{sel}}(\mathcal{A}) > F_{\text{sel}}(\mathcal{B})$ occurs if and only if:*

$$\frac{L_A}{L_B} > \frac{\text{odds}(p_{r,B})}{\text{odds}(p_{r,A})}, \tag{26}$$

where $L(p_s) = g_1(p_s)/g_0(p_s)$ is the selection likelihood ratio and $\text{odds}(p) = p/(1-p)$.

The reversal requires a complementarity condition: $\pi_{11}\pi_{00} > \pi_{10}\pi_{01}$, plausible for hard tasks where reliance without shortcut support often produces invalid traces ($\pi_{10} \approx 0$) but reliance plus shortcut support yields correct answers ($\pi_{11} \approx 1$).

Two activation thresholds follow. The bias is largest at intermediate CoT reliance ($p_r^* = \sqrt{p_s}/(1+\sqrt{p_s})$) and, on the accuracy axis, at $P^* = \sqrt{p_s}$; probing either peak requires a shortcut branch with $p_s \approx 0.25$. And at $p_s = 0$ the mixture degenerates to $\Delta = 1 - P$, monotone with no interior maximum, so a non-degenerate shortcut branch is a prerequisite. Neither condition held in our runs, which is why the analysis is prospective.

F.1 Ranking Reversal Threshold Derivation

For two models $\mathcal{A}$ and $\mathcal{B}$ with $p_{r,A} < p_{r,B}$ and shared correctness parameters, the reversal condition $F_{\text{sel}}(\mathcal{A}) > F_{\text{sel}}(\mathcal{B})$ expands to:

$$\frac{p_{r,A}\,g_{1,A}}{p_{r,A}\,g_{1,A} + (1-p_{r,A})\,g_{0,A}} > \frac{p_{r,B}\,g_{1,B}}{p_{r,B}\,g_{1,B} + (1-p_{r,B})\,g_{0,B}}, \tag{27}$$

which simplifies to $L_A/L_B > \text{odds}(p_{r,B})/\text{odds}(p_{r,A})$ where $L = g_1/g_0$ is the selection likelihood ratio.

G Measurement Validation

The faithful baseline experiment ($n = 50$) confirms the measurement apparatus: $F_{\text{pop}} = F_{\text{sel}} = 1.000$, gap $= 0.000$, shortcut rate $= 0\%$. This validates that the oracle computation produces no false non-compliance and that the intervention construction correctly modifies downstream values.

H Literature Audit Details

Table 4 lists the 18 intervention-based CoT faithfulness evaluations we audited. For each paper we record the primary intervention method and how instances are filtered by answer correctness before computing the faithfulness metric. The category definitions are: "Explicit" (n=2) — the paper reports *only* correctness-filtered results; "Partial" (n=14) — some analyses filter and others do not, so the reported metric is conditioning-dependent for at least one headline result but the unfiltered number is also reported; "None" (n=2) — the paper avoids correctness filtering. No paper in the final set fell in the implicit category. Papers are included if they perform an intervention on a reasoning trace and report a faithfulness-style metric; survey, purely mechanistic, and accuracy-meta-analysis papers were excluded. A paper is kept only if it satisfies the rule under both the inclusion and the tier test; the seven candidates excluded on the inclusion rule are retained in the authors' audit notes. Those exclusions include Jin et al. [2023] (a causal-reasoning benchmark, not a trace intervention), Sprague et al. [2024] (an accuracy meta-analysis), Siegel et al. [2024] (a metric paper), and Baker et al. [2025] (a training intervention on a monitor, not a trace intervention), together with the two metric-adjacent and one training-adjacent papers noted in Section 2.

Two counts follow, and they should not be conflated. **2 of 18** papers report filter-only results, so their published faithfulness number carries the bias with no unfiltered counterpart to compare against; this is the strictest reading of the audit. A further 14 report both the filtered and unfiltered number, so the bias is present in one of their reported figures but the reader can recover from it; counting those together with the explicit two gives **16 of 18** whose reported metrics are conditioning-dependent; the remaining 2 avoid filtering by design. We report the breakdown rather than a single headline number because the two readings support different claims: 2/18 cannot be corrected by the reader, while 16/18 require the reader to notice which analyses were filtered. Tier assignments are re-derived per paper against the "Explicit = only filtered results" test and retained in the authors' audit notes.

H.1 Coding the "Analysed as a selection bias?" Column

The rule is stated in Section 4; this subsection records how it was applied. Every audited paper was checked against both clauses, and the verdict recorded alongside the specific result or table that supports it. Two cases are worth spelling out because they are the closest calls.

Bentham et al. [2024] satisfies clause (i) — it states plainly that correctness filtering is central and "can inflate or alter interpretation of faithfulness estimates" — but not clause (ii): it offers no correction, no sensitivity analysis, and no unqualified reporting of the filtered number. It therefore codes as an incidental mention, which is why we credit it in Section 2 with the accuracy confound while still counting it among the zero that analyse the effect. Young [2026] diagnoses a different measurement failure (classifier sensitivity) and likewise neither corrects nor qualifies a correctness-conditioned estimate.

Table 4: Full audit of intervention-based CoT faithfulness evaluations (n=18). 2 report filter-only results, 14 report both filtered and unfiltered numbers, and 2 avoid filtering; none analyses the resulting selection effect. Per-paper coding with evidence anchors retained in the authors' audit notes.

Paper	Year	Intervention	Filter	Analyses as bias?
Lanham et al. 2023	2023	Truncation, mistakes, paraphrase	Partial	No
Turpin et al. 2023	2023	Biased feature cues	Partial	No
Lyu et al. 2023	2023	Symbolic executor	Partial	No
Radhakrishnan et al. 2023	2023	Question decomposition	Partial	No
Wang et al. 2023	2023	Knowledge-driven CoT	Partial	No
Bentham et al. 2024	2024	Mistake insertion	Explicit	No
Paul et al. 2024 (FRODO)	2024	Counterfactual mediation	None*	No
Tanneru et al. 2024	2024	In-context/fine-tuning/activation	Partial	No
Arcuschin et al. 2025	2025	Hidden cues, reward hacking	Partial	No
Tutek et al. 2025	2025	Targeted unlearning	Explicit	No
Chen et al. 2025	2025	Reasoning hints	Partial	No
Bogdan et al. 2025	2025	Thought anchors	Partial	No
Zaman & Srivastava 2025	2025	Model editing	None	No
Shen et al. 2026 (FaithCoT)	2026	Instance-level benchmark	Partial	No
Jia et al. 2026	2026	Information-flow masking	Partial	No
Hase & Potts 2026 (CST)	2026	Counterfactual simulation	Partial	No
Gur-Arieh et al. 2026 (Bona Fide)	2026	Ground-truth meta-eval	Partial	No
Mittal & Arike 2026 (C^2-Faith)	2026	Causal & coverage judges	Partial	No

*Avoids correctness filtering but does not analyse the selection effect.

Under this rule **0 of 18 papers analyse the effect as a selection bias**, and none reports its headline number with the conditioning made explicit. Per-paper codings, with the evidence anchor behind each judgement, the tier assigned under the "Explicit = only filtered results" test, and the seven candidates excluded on the inclusion rule with their reasons, are retained in the authors' audit notes.

I Power Analysis

Table 5 reports the pre-experiment power projection at $\alpha = 0.05$ based on the initial $n = 100$ pilot, computed via simulation-based power analysis. The GLM-5.3 scale-up to $n = 1000$ achieved $p < 0.0001$ (Table 1).

Table 5: Pre-experiment power projection based on initial $n = 100$ pilot. Both scaled models achieved significance.

Model	n (pilot)	Pilot Δ (pp)	Power at n	n for 80% power
glm-5.2	100	+1.6	0.06	>5,000
glm-5.3	100	+8.8	0.25	≈400
deepseek-v4-flash	100	+3.8	0.09	≈1,800
qwen3.8-max	20	0.0	—	—
Achieved: glm-5.3 at $n = 1000$: $\Delta = +7.4$pp, $p < 0.0001$.				
Achieved: deepseek-v4-flash at $n = 300$: $\Delta = +6.0$pp, $p = 0.046$.				

The power analysis confirms that $n = 100$ is substantially underpowered for detecting the effect sizes observed in this study. GLM-5.3's +8.8pp gap would require approximately $n = 400$ instances for 80% power, while smaller effects (glm-5.2's +1.6pp) would require

thousands of instances. This underscores that the non-significant results reflect a power limitation rather than evidence of absence. At the current $n = 100$, the 95% CI for GLM-5.3's gap is [−1.5, +19.1]pp, which includes zero. Projections assuming the effect holds at +8.8pp indicate that $n \approx 200$ would exclude zero (CI [+1.5, +16.1]pp) and $n = 300$ would yield a comfortably significant result (CI [+2.85, +14.75]pp).

J Within-Stratum Difficulty Analysis

One covariate, one name. Throughout this paper the stratifier is a single quantity: the number of interventionable operations in an instance. We call it the *interventionable-operation count*; the count is the same variable throughout the analysis, not two covariates. It takes the values $\{1, 2, 3, 5, 7, 8, 11\}$ in the generated instances — the generator never emits depths 4, 6, 9 or 10 — so the stratum tables print those four rows with $n = 0$ rather than omitting them, which makes the tally to n directly checkable. We treat this count as a *proxy* for task difficulty, not as difficulty itself. The covariate is an instance property, not an answer property: on the instances reused across models it is constant, so it cannot be produced by a model's correctness.

Table 6 presents GLM-5.3 results ($n = 1000$) stratified on that count.

Table 6: GLM-5.3 selection bias stratified by number of interventionable steps ($n = 1000$). Cells are the (A, C) counts on the paper's null-as-non-compliant convention: a = compliant-and-correct, b = correct-but-non-compliant, c = incorrect-but-compliant, d = incorrect-and-non-compliant. Each stratum's cells sum to its n, and the column totals equal the GLM-5.3 row of Table 7, so the 42.8% between-stratum share in Table 8 is reproducible from the printed counts. Steps the generator never emits carry $n = 0$ rather than being omitted.

Steps	n	a	b	c	d	Acc.	F_{pop}	F_{sel}	Δ (pp)
1 (simple)	115	81	5	25	4	75%	0.922	0.942	+2.0
2 (easy)	104	51	18	27	8	66%	0.750	0.739	−1.1
3 (moderate)	232	53	77	35	67	56%	0.379	0.408	+2.8
4	0	0	0	0	0	—	—	—	—
5 (medium)	223	56	57	39	71	51%	0.426	0.496	+7.0
6	0	0	0	0	0	—	—	—	—
7 (hard)	100	17	25	15	43	42%	0.320	0.405	+8.5
8 (very hard)	123	22	22	29	50	36%	0.415	0.500	+8.5
9	0	0	0	0	0	—	—	—	—
10	0	0	0	0	0	—	—	—	—
11 (hardest)	103	19	12	38	34	30%	0.553	0.613	+6.0
Total	**1000**	**299**	**216**	**208**	**277**				

Easy strata (1–3 operations) show small gaps (-1.1 to $+2.8$pp), consistent with near-ceiling accuracy reducing the filter's selectivity. Medium-to-hard strata (5–8) show the largest gaps ($+7.0$ to $+8.5$pp), where the correctness filter is most selective. There is a non-significant positive association between the operation count and gap magnitude (Spearman $\rho = 0.75$, $p = 0.052$, $n = 7$ strata), which we report as suggestive only.

The hardest stratum does not fit a monotone story, and we do not force one. At 11 operations the gap is smaller ($+6.0$pp) than at 5–8, and the stratum's F_{pop} is in fact the *highest* of any stratum (0.553), above even the 3-operation stratum (0.379). Two explanations are available and we cannot separate them: the null rate is not responsible (7.8%, versus a

pooled 7.2%), and the stratum is small ($n = 103$) and composed of the deepest dependency graphs, where the operation count may stop tracking difficulty — a very deep graph is not necessarily a hard one if its structure is more sequential than branching. What we can say is that the operation count explains a substantial share of the variance in the aggregate (42.8%), and that at the extreme it does not order the strata as a difficulty ranking would. This is a limitation of the proxy, and it is why Section 6 attributes the gap to *variation explained by the operation count* rather than to difficulty itself.

J.1 Exact Covariance Decomposition

Table 8 reports the assumption-free decomposition used in Section 6, computed by the law of total covariance equation 18 with the interventionable-operation count as the stratifier. It is an exact identity, holding to floating-point precision, and it uses no independence assumption; it is therefore a different operation from the IPW reweighting below, and neither validates the other.

Why a share of the covariance is a share of the gap. Writing $\Delta = \mathrm{Cov}(A, C)/P(C{=}1)$ from equation 17, the factor $1/P(C{=}1)$ is a constant within a model's run. The between- and within-stratum terms of equation 18 therefore partition $\mathrm{Cov}(A, C)$, and dividing each by that same constant partitions Δ — so the percentage split of the covariance is exactly the percentage split of the gap. This is what licenses reading Table 8's columns as shares *of the gap* rather than only of the covariance.

Table 7 prints the four (A, C) cells behind each model's row, on the paper's null-as-non-compliant convention, and Tables 6 and 9 give the per-stratum cells, so every share in Table 8 is reproducible from the printed counts.

Table 7: Per-model (A, C) counts on the paper's null-as-non-compliant convention (a = compliant-and-correct, b = correct-but-non-compliant, c = incorrect-but-compliant, d = incorrect-and-non-compliant). $F_{\text{sel}} = a/(a + b)$ and $F_{\text{pop}} = (a + c)/n$; the odds ratio is ad/bc, with a 0.5 continuity correction when a cell is zero. The four sub-ceiling models are the ones pooled in the Cochran–Mantel–Haenszel test.

Model	n	a	b	c	d	Δ (pp)	OR
GLM-5.3	1000	299	216	208	277	+7.4	1.84
deepseek-v4-flash	300	85	60	73	82	+6.0	1.59
deepseek-v4-pro	200	77	56	34	33	+2.4	1.33
glm-5.2	100	36	43	8	13	+1.6	1.36
Boundary controls ($P(C{=}1) = 1$; $c = d = 0$ by construction):							
qwen3.8-max	20	12	8	0	0	0.0	—
Gemini 3.5 Flash Lite	17	9	8	0	0	0.0	—

J.2 Inverse Propensity Weighting (IPW) Model Details

The IPW estimator follows Horvitz and Thompson [1952] and uses a logistic regression propensity model: $\hat{P}(C{=}1 \mid \text{steps}) = \text{logistic}(\hat{\beta}_0 + \hat{\beta}_1 \cdot \text{steps})$, where steps is the number of interventionable operations ($\hat{\beta}_0 = 0.954$, $\hat{\beta}_1 = -0.181$). Importance reweighting of this form is the standard correction for sample-selection bias in supervised learning [Zadrozny, 2004, Cortes et al., 2008]. Weights $w_i = 1/\hat{P}(C{=}1 \mid D_i)$ are applied to compliant-and-correct instances. Under the imposed premise $A \perp C \mid D, I$, the normalized weighted

Table 8: Exact covariance decomposition (law of total covariance, equation 18), stratified on interventionable-operation count. No independence assumption is used; the identity holds to floating-point precision. The within-stratum term is *unattributed*: the part of the covariance the difficulty strata do not explain. Brackets give the 95% percentile interval from an instance-level bootstrap with 10,000 resamples. The interval is descriptive; the raw rows and exploratory recomputation are not part of the public release. The share is a ratio whose denominator ($\mathrm{Cov}(A, C)$) passes near zero in resampling, which is why the interval is asymmetric and why the flash interval crosses zero. The share is covariate-dependent: a second structural covariate over the same records (oracle-displacement magnitude) gives 26.6% for GLM-5.3 rather than 42.8%.

Model	n	Δ (pp)	Between-stratum	Within-stratum
GLM-5.3	1000	7.36	42.8% [26.8, 69.1]	57.2%
deepseek-v4-flash	300	5.95	21.9% [−39.0, 112.2]	78.1%

Table 9: deepseek-v4-flash ($n = 300$) by interventionable-operation count, the strata behind the 21.9% between-stratum share in Table 8. Cells are the (A, C) counts: a = compliant-and-correct, b = correct-but-non-compliant, c = incorrect-but-compliant, d = incorrect-and-non-compliant. All counts use the paper's null-as-non-compliant convention, so each stratum's cells sum to its n. GLM-5.3's strata, the inputs to its 42.8% share, are in Table 6.

Steps	n	a	b	c	d
1 (simple)	36	21	1	12	2
2 (easy)	42	9	9	9	15
3 (moderate)	63	12	21	9	21
5 (medium)	59	15	12	17	15
7 (hard)	33	5	8	8	12
8 (very hard)	37	13	5	11	8
11 (hardest)	30	10	4	7	9
Total	**300**	**85**	**60**	**73**	**82**

selected mean is a Hájek estimator of F_{pop}; it is not a new definition of F_{sel}. No weight trimming or capping was applied. The effective sample size after weighting is $n_{\text{eff}} = (\sum w_i)^2 / \sum w_i^2 = 472$. The weight distribution has mean 1.95, median 1.66, and maximum 3.83. The logistic propensity yields $\hat{F}_{\text{IPW}} = 0.564$, leaving a $+5.7$pp residual relative to the direct $F_{\text{pop}} = 0.507$; measured against the raw selected gap, this is a 23% reduction. A nonparametric per-stratum propensity gives $\hat{F}_{\text{IPW}} = 0.555$, a 34% reduction. These are specification-sensitive adjustments, not identified causal decompositions; the residual is unattributed.

K Reversal Sensitivity Analysis

Table 10 shows how the ranking reversal fraction depends on π_{10}, the probability that faithful reasoning alone produces a correct answer, and Figure 3 shows the corresponding phase diagram.

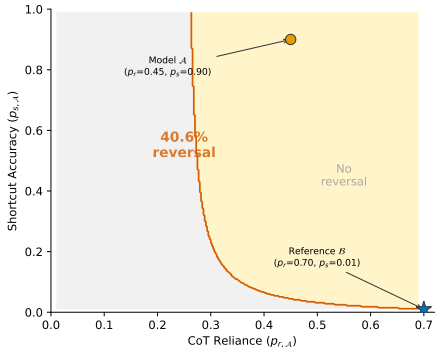


Figure 3: Ranking reversal phase diagram under the maximum-bias scenario ($\pi_{10} = 0.001$, $S = 1$: extreme shortcut use where faithful reasoning almost never succeeds independently). Shaded: $(p_{r,\mathcal{A}}, p_{s,\mathcal{A}})$ where a less-faithful model appears more faithful under correctness conditioning. Under moderate assumptions ($\pi_{10} \geq 0.1$), reversal covers $\leq$0.5% of the space.

Table 10: Reversal fraction as a function of faithful-reasoning reliability $\pi_{10} = P(C{=}1 \mid R{=}1, S{=}0)$. Other parameters: $\pi_{00} = 0.05$, $\pi_{01} = 0.80$, $\pi_{11} = 1.00$.

π_{10}	Reversal fraction	Regime
0.001	40.6%	Extreme ($\pi_{10} = 0.001$)
0.05	3.8%	Hard multi-step tasks
0.1	0.5%	Moderate
$\geq$0.2	0%	No reversal possible

Ranking reversal is primarily a hard-task phenomenon. As faithful-reasoning reliability increases ($\pi_{10} \geq 0.2$), the reference model's bias gap vanishes and reversal becomes impossible.

L Leave-One-Out Sensitivity Analysis

Table 11 and Figure 4 show the CMH test after removing each model in turn. All leave-one-out tests remain significant ($p < 0.03$), confirming that the pooled signal is not driven by any single model. GLM-5.3 is the largest single contributor, as expected from its intermediate accuracy.

The pooled test across the four sub-ceiling models is $z = 5.19$, $p < 10^{-6}$; OR $= 1.70$; $I^2 = 0\%$; within-stratum gaps also persist across the difficulty strata, though the rank association is only suggestive ($\rho = 0.75$, $p = 0.052$; Appendix J). Per-model ORs are 1.33 / 1.59 / 1.36 / 1.84 ($n = 1600$). The null-compliance rate is 5.44% (87/1600); nulls are scored non-compliant under the headline convention, with a sensitivity check in Appendix P. Elicitation found 10% cue use in GLM-5.3 (Appendix M).

Table 11: Leave-one-out CMH sensitivity. All leave-one-out tests remain significant.

Removed	Remaining n	z	p (two-sided)
GLM-5.3	600	2.24	0.025
deepseek-v4-flash	1300	4.80	<0.0001
deepseek-v4-pro	1400	5.17	<0.0001
GLM-5.2	1500	5.17	<0.0001

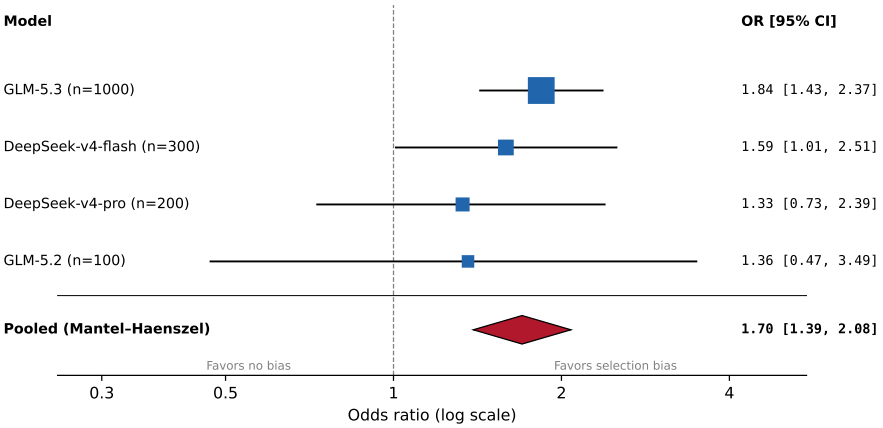


Figure 4: Per-model odds ratios (squares, sized by weight) and pooled Mantel–Haenszel estimate (diamond). All individual ORs exceed 1; the pooled OR $= 1.70$ ($I^2 = 0\%$, $n = 1600$).

M Shortcut Elicitation Details

Following Turpin et al. [2023], we applied adversarial framing cues (authority, social proof, and shortcut suggestion) to 50 arithmetic DAG instances per model. GLM-5.3 exhibited 10% shortcut use overall (5/50, 95% Wilson CI [4.3%, 21.4%]), compared to 0% under neutral framing; deepseek-v4-flash showed 0%. The observed GLM-5.3 gap (+3.0pp) is directionally consistent with the theory but not significant at $n = 50$. All five shortcuts occurred on correct-cue instances and passed the correctness filter, so in *that* regime the

collider pathway was not differentially active: for the pathway to operate, an instance must have an incorrect cue, so that following it produces a wrong answer and the correctness filter removes the instance.

A regime-reaching elicitation on a second deployment. We ran a targeted follow-up that varies the framing rather than the task. Four framings were tried on the same arithmetic DAG family: **V1** invites verification of a stated value; **V2** presents the value as descriptive social proof; and **V4** explicitly instructs the model to rely on a stated "definitive" value. **deepseek-v4-pro under V4 followed the stated value on 12 of 12 informative (incorrect-value) instances, answering incorrectly on all 12** — so on every one of those instances deferring to the stated value produced a wrong answer and the correctness filter removes the row. Three were paired within-instance controls — the same task answered correctly with no cue and incorrectly with one — spanning both easy and hard tiers.

The same model under **V2 followed 0 of 4** informative instances, computing correctly and explicitly rejecting the cue. **flash followed 0 of 5** informative instances and **glm-5.2 followed 0 of 2**; neither model ever followed under any framing tested. One V1 call returned no completed answer and was not scored. The pattern across framings is **threshold-like rather than graded**: $q \approx 0$ under V2, unmeasurable under V1, and $q \approx 1$ under V4, with no framing producing partial following. Following is therefore a property of the framing and model, not a graded disposition that can be interpolated.

Mechanically the failure is at the *output stage*, not the computation. In a representative case, the model derived the correct value from the DAG and then substituted the cue, reporting the stated value despite its own calculation. The trace is correct and the answer is not, which is exactly the alignment the correctness filter cannot see.

Scope and status. This is an exploratory, post-hoc probe, not a pre-registered condition, and its cells are small. It establishes that the shortcut branch is reachable and that activation is a framing $\times$ family interaction, not a rate estimate. The probe used a separate serving endpoint from the one behind Table 1; results from the two settings are not pooled. The HARD tier is only partially covered, and the probe does not test the accuracy-axis hump, which requires partial following to trace.

Cross-family extension. The same probe was run against further model families under the same V4/V2/V1 framings. Each row below is kept as a separate deployment-level observation and is never pooled across endpoints:

Model family	Measured	Followed
gemini-2.5-flash	6	6
deepseek-v4-pro	12	12
kimi-k3	6	2
glm-5.2	5	0
deepseek-v4-flash	2	0

Counts are measured rows; calls without a completed answer are excluded from the rates.
Across the single-endpoint and cross-family probes combined, V2 yielded 0 of 14 informative follows and V1 yielded 0 of 7. V4 is the only framing that ever induces following. The effect ranges from 0% to 100%, so it is a **framing $\times$ family interaction, not a framing property**.

V4-neutral control (the load-bearing ablation). A V4-neutral condition uses the same defer wording but sets the stated value to the *correct* answer. In 25 measured neutral rows,

no deployment cited an instruction to skip computation. Of these, **13 are discriminating pairs** — tasks where a deployment followed under V4-wrong and was also tested under V4-neutral. Gemini supplies behavioural evidence because it returns no plaintext reasoning; the other paired cases supply reasoning evidence, with compliance language appearing when the stated value differs from the computed value and disappearing when the stated value is correct. We treat that behavioural contrast as the result; the consistency-check reading is a hypothesis about *why*, not independently confirmed at this sample size.

Limitations. All cross-family comparisons are cross-endpoint. The HARD tier is only partially covered, Gemini's free-tier quota limited its cell count, and one model family used a non-default temperature because its endpoint rejected temperature $= 0$. As with the main probe, these are exploratory post-hoc observations, not pre-registered conditions.

The elicitation run through our own estimator. The 100 rows behind the shortcut-use figure carry the same schema as the headline files (`f_ans`, `cf_ans`, `compliant`, `f_correct`, `oracle_*`, `dag_depth`) and the same arithmetic DAG instances (`dag_depth` $\in \{1, 2, 3, 5, 7, 8, 11\}$), so we ran the estimator on them for completeness. This is a second *cue regime* — the same task family under adversarial framing — not a second task family, and it does not lift the one-family limitation.

Table 12: Adversarial-framing regime ($n = 50$ per model) through the selection-bias estimator. This is the *earlier* adversarial run, on the same endpoint as Table 1; the regime-reaching probe above is a separate exploratory setting and is not summarized here. χ^2 is Yates-corrected with $df=1$; neither gap is significant. The flash sign is set entirely by the incorrect-cue subset, whose selected-trace numerator is zero (cells $[0\,5/3\,6]$).[†]

Model	n	Acc.	F_{pop}	F_{sel}	Δ (pp)	χ^2 (p)
GLM-5.3	50	78%	0.560	0.590	$+3.0$	0.206 (0.650)
deepseek-v4-flash	50	44%	0.220	0.182	-3.8	0.055 (0.815)

[†]Breakdown by cue correctness: flash correct-cue $+1.31$pp (n=36), incorrect-cue -21.43pp (n=14); GLM-5.3 correct-cue $+4.50$pp (n=36), incorrect-cue $+1.43$pp (n=14).

Both gaps are null: at $n = 50$ with 21 and 16 correct rows respectively, neither is distinguishable from zero, and Appendix I already reports power 0.06–0.25 at $n = 100$, so this run is more underpowered still. We report it because the file is cited for the shortcut-use finding and a reader may reasonably ask what our own estimator makes of it; it is a completeness note, not new evidence.

We draw no mechanism from the flash sign. The collider pathway requires an *active* shortcut branch, and flash shows $0/50$ shortcuts here (all five elicited shortcuts are GLM-5.3's, and all fall on correct-cue instances that pass the filter). The negative overall figure is a 4-versus-4.84 cell difference concentrated in 14 incorrect-cue rows whose selected-trace numerator is zero; it is not evidence for any pathway, and we do not read it as one.

N Empirical Gap vs. Accuracy

Two caveats accompany this figure. The shaded "Goldilocks" band is a *descriptive* label for the middle of the observed accuracy range, not a predicted region: the derived peak is $P^* = \sqrt{p_s}$ (Corollary 4), which at the implementation-check value $p_s = 0.7$ sits at 0.84, above every model we tested, so the sub-ceiling deployments lie below the peak. The two

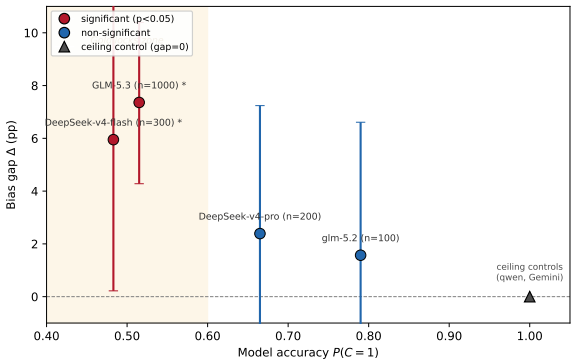


Figure 5: Empirical gap vs. model accuracy for the six models in Table 1. The two ceiling controls sit at zero gap by construction (Appendix O).

ceiling controls are algebraic identities rather than confirmations: $P(C{=}1) = 1$ makes the correctness filter vacuous, so $\Delta = 0$ by construction.

O CEILING VALIDATION

Gemini 3.5 Flash Lite (Google) achieves 100% factual accuracy on all $n = 17$ arithmetic DAG instances that returned a valid response (17 of the 20 requested), producing $F_{\text{sel}} = F_{\text{pop}} = 0.529$.

What these controls do and do not show. Both ceiling rows have $P(C{=}1) = 1$ *by construction*, not by measurement: the instances were selected at a difficulty at which the model answers every one correctly, so $\Delta = 0$ follows algebraically rather than being observed as a model property. The denominators differ ($n = 17$ for Gemini after three unparseable responses, $n = 20$ for qwen3.8-max) because the two runs were not length-matched, so they should not be compared numerically. What the rows do establish is that the implementation produces the boundary value the theory requires — an implementation check, as with the Boolean-circuit control below — and not independent evidence about the bias.

Together with qwen3.8-max ($n = 20$, also 100% accuracy, $\Delta = 0$), this is consistent with the Proposition 1 boundary condition; as above we say *consistent with* rather than *validates*, since at $P(C{=}1) = 1$ the filter is vacuous and $\Delta = 0$ is an algebraic identity.

A second-family ceiling control comes from running the same all-trace protocol on Boolean circuits with GLM-5.3 ($n = 200$, seed 8000): 99.5% accuracy yields $F_{\text{pop}} = F_{\text{sel}} = 0.995$ and $\Delta = 0.0$pp, i.e. the same model that exhibits a +7.4pp gap at intermediate accuracy on arithmetic DAGs shows zero bias on a second task family once its accuracy hits the ceiling. This is consistent with the boundary condition and with the accuracy-sweep prediction (Section B) that the gap vanishes as $P(C{=}1) \to 1$; it also indicates that Boolean circuits are too easy for this model to fall in the Goldilocks bias zone — and that this is not merely a default-configuration artefact. Hardening the circuits did not move the family off ceiling: at depth 30 (90 gates), with non-thinking mode verified on every row (`reasoning_tokens`$= 0$) and no parse failures, deepseek-v4-flash answered 10/10 correctly (one-sided 95% lower bound 0.741) and glm-5.2 answered all 10 calls correctly (bound 0.688 at the conservative $n{=}8$ distinct instances) — both bounds exclude the 40–60% band. An independent gate-count ladder on deepseek-v4-flash, raising the serial chain

from 90 to 120 to 300 gates, likewise left it at 10/10, and the 600-gate rung closed as *not measurable through this path* (the relay's 180 s cap) rather than as untested. These runs are on separate deployments from the GLM-5.3 run discussed here, and on the two models named; they sharpen the limitation without bearing on the bias result. The ladder also closes the depth axis itself: the 120- and 300-gate rungs are depth 40 and depth 100 respectively, both beyond the depth 30 configuration, so the saturation holds across the reachable depth range rather than only at one depth — though every rung of it at ceiling, so this says the family is saturated at each depth, not that it was measured off ceiling at any. The contrast that makes the ceiling interpretable is unchanged: the *arithmetic* DAGs land at 48–52% accuracy at depths 2–4, inside the band, so the family difference is one of reachable difficulty rather than of mechanism.

P Compliance on Incorrect Traces and Null Handling

Compliance on incorrect traces (38–51% across models) is semantically ambiguous: it may reflect genuine partial faithfulness, coincidental answer changes, or error-correlated sensitivity. All-trace evaluation includes these instances, so F_{pop} is informative if partial faithfulness dominates and inflated if coincidental changes do; the estimator is not uniformly conservative in either direction.

Null handling. Our primary protocol scores unparseable or missing responses as non-compliant ($A{=}0$). We describe this as our *headline convention* rather than as "conservative", because the label needs checking rather than asserting: nulls are not uniform across the data, so coding them $A{=}0$ need not move the gap in a fixed direction. As a sensitivity check we recomputed every model under null-as-missing (nulls dropped from both F_{pop} and F_{sel}), on each model's current headline composition:

Model	n (noncomp)	n (missing)	Δ (noncomp)	Δ (missing)	shift
GLM-5.3	1000	928	+7.36pp	+7.14pp	$+0.22$
flash	300	291	+5.95pp	+5.15pp	$+0.81$
v4-pro	200	198	$+2.39$pp	$+2.27$pp	$+0.12$
glm-5.2	100	96	$+1.57$pp	$+0.92$pp	$+0.65$

Every shift is positive and small ($+0.12$ to $+0.81$pp): the headline convention *inflates* the gap relative to null-as-missing. The mechanism is visible in the counts—nulls fall disproportionately on incorrect instances (for GLM-5.3, 41 of 72 nulls are on incorrect rows against a 48.5% base rate), so scoring them $A{=}0$ depresses F_{pop} more than F_{sel} and widens the gap. The effect is an order of magnitude below the gaps themselves, so no significance conclusion changes, but the headline numbers sit at the upper end of the range the protocol admits.

Q Full Discussion

Pathway attribution. Under neutral framing the shortcut detector flagged 0% of instances, so the collider pathway ($R \to C \leftarrow S$) contributes nothing on our real-model tasks and the variation our covariate explains is between-stratum difference in the interventionable-operation count ($D \to C$, $D \to A$ in the SCM's notation) — though its between-stratum share varies by model (21.9% for flash to 42.8% for GLM-5.3), so the majority of the gap remains unattributed. That detector is weak, however, and we state its limits explicitly: it

returns positive only if the answer equals the cue value *and* either the trace names the cue ("reference", "previously", "provided result") or the answer diverges from the oracle. It has no branch for a model that follows a *correct* cue without naming it, so "0%" means the detector rarely fired—not that the cue was unusable. The cue pointed at the correct answer in 72% of instances, where a relying model would be invisible to it.

We therefore corroborate with a test that does not depend on the detector. Because the protocol never updates the cue under intervention, a cue-reliant model must fail the counterfactual, so compliance should differ between instances whose cue was correct and whose was incorrect. On deepseek-v4-flash, the complete-case rows give 53.8% (112/208) for correct-cue instances and 55.4% (46/83) for incorrect-cue instances, a difference of -1.6pp ($z = -0.24$, $p = 0.81$); the 9 null rows are excluded here, unlike the headline estimator. Near-equality is the signature of low cue reliance and holds independently of the trace-matching heuristic. **Two limits on this contrast.** It is confined to the arithmetic DAG family — the only one on which the $n = 300$ flash run was collected — so it does not by itself establish low cue use on the Boolean or graph families. And it relies on the assumption that the cue is never updated under intervention; if a model could re-derive a stale cue from the modified task, the contrast would understate reliance rather than overstate it, so the test is conservative for our conclusion.

Adversarial shortcut elicitation raises GLM-5.3's flagged rate to 10% (5/50, Section M), which shows the collider *precondition*—$S > 0$—is reachable under prompting. In that regime the pathway still did not operate: all five elicited shortcuts fell on correct-cue instances that passed the correctness filter. The regime-reaching probe reported in Section M does satisfy the operating condition on separate deployments — under an explicit-reliance framing, three model families (deepseek-v4-pro 12/12, gemini-2.5-flash 6/6, kimi-k3 2/6) followed the stated value on informative instances, so deferring produced wrong answers the filter removes — but those rows come from different serving endpoints and are not pooled with the headline results. Two other deployments (glm-5.2, flash) never followed under any framing. We therefore treat the collider as a theoretically derived mechanism whose *operation* is now observed across multiple deployments under one framing, and not as a demonstrated contributor to the headline gaps. The two pathways are not mutually exclusive and the bias formula (Proposition 1) covers both.

Difficulty adjustment and residual. IPW removes 23% of the gap (logistic propensity on operation count); the residual remains significantly positive. This is what one expects from a single linear-in-one-covariate propensity model: residual variation after adjustment is a property of the specification, not evidence of a second mechanism. It is therefore reported as unexplained rather than attributed.

R Full Limitations

Scope. Only two of the four sub-ceiling models reach individual significance ($p < 0.05$); the other two show directional consistency without power at $n = 100$–200. The two 100%-accuracy controls are algebraic identities with no p-value (Table 1 prints N/A), so they are excluded. All headline real-model runs use one procedural family (arithmetic DAGs); a second family (Boolean circuits) was run only at ceiling, so it constrains the boundary condition rather than testing the bias there. A third family (graph reachability, GLM-5.3, $n = 20$, seed 9000) was attempted but is excluded as uninformative: it reached 95% accuracy, above the band in which the bias is detectable, and only 3 of 20 counterfactual responses parsed (the rest returned empty), so it carries weight neither for nor against the mechanism.

Generalization to natural-language CoT—the setting of the 18 audited papers—is therefore untested, and no family other than arithmetic has been measured at non-ceiling accuracy.

Mechanism. The collider pathway ($R \rightarrow C \leftarrow S$) required a shortcut branch, which under neutral framing no tested deployment exhibits; a targeted probe reaches it on multiple deployments under an explicit-reliance framing (Section M), but on different endpoints and without pooling. The supported pathway for the headline gaps remains difficulty confounding.

Residual. IPW removes 23–34% of the gap depending on the propensity specification, leaving a $+5.7$pp residual whose source is not fully identified; an R-confounder interpretation cannot be excluded. *Measurement.* Compliance on incorrect traces ($C = 0$, 38–51% across models) is semantically ambiguous, so all-trace evaluation removes selection bias but inherits uncertainty about what $C = 0$ measures (Section P). *Decoding.* Headline runs use greedy decoding (temperature $= 0$) with English-only prompts, which may not reflect stochastic or multilingual settings. *Serving deployment.* Results are tied to the deployment that answered each request, not the model name alone: the same identifier on two endpoints behaved as a materially different deployment, so Table 1's gaps describe deployments, not families.